
Pattern Arithmetic: Typed Algebra of Bounded Units on a Phase-Colored Sphere

Sandeep Lakshminarayan Chiluveru

October 2, 2026

Abstract

Ordinary arithmetic treats the integer 1 as a unit of truth: $1 = 1$ and $1 + 1 = 2$. Counting is only one algebra. Colors and musical notes are also units of truth; combining them yields a hue or a chord, not a larger integer. This paper treats a pattern as a higher-order unit and represents one pattern by one typed point $x = (r, \hat{n}, \chi, \tau)$ on a unit sphere $\mathcal{S} \subset \mathbb{R}^3$ ¹: strength $r \in [0, 1]$, direction $\hat{n} \in S^2$ (equivalently (θ, ϕ)), internal phase-color $\chi \in [0, 2\pi)$, and sort τ . Pattern Arithmetic is the algebra of such units. Five verbs answer five questions. Bundle ($\uplus$) lists who is present and does not double identical states ($A \uplus A = A$); a Bundled unit may be removed ($\{A\} \setminus A = \emptyset$). MixMerge ($\boxplus$) reads a composite; anti-phase cancellation to black is that reading, not a Bundle. Transform (T) walks an order so the same ingredients may form different expressions. Offset reuses a walk. Assembly ($\boxtimes$) relates two units without fusing them — the first operator here that needs anything outside the 4-tuple: an external seat. Units left alone descend an energy until a stated threshold freezes them, resting in one camp or several; left longer, all would merge into one. Bundle $\rightarrow T \rightarrow$ result is one pipeline; assemble $\rightarrow$ rest is the other. Division as unmixing is shown ill-posed for this algebra, not attempted here. Words and numbers occupy the same sphere under different encodings and meet in a document by order and typed action, not by blending. The aim is a single bounded register that carries overlapping identities, structural paths, and mixed prose without a discrete grid. The sphere and its operators are one proposed encoding of a narrower claim — a combination rule should match the unit — not a theory of language, physics, biology, or computation.

¹Three dimensions is the standing assumption throughout — every axiom, proof, and worked example below uses $\mathbb{R}^3$ unless a section states otherwise. Section 3.4 extends T to $S^{N_{\text{dim}}-1}$ for $N_{\text{dim}} \neq 3$; that is a named exception, not a change to the default, as is the general-dimension form of offset’s transport in §4.1.

1 Introduction

Ordinary arithmetic uses the integer 1 as a unit of truth. The identity $1 = 1$ says that two tokens are of the same kind. The rule $1 + 1 = 2$ says that combining two such tokens produces a larger quantitative state. That algebra answers the question *how many?*

It is not the only algebra. A color and a musical note are also units of truth, but their combination does not enlarge a pile. Under additive light, red and green yield yellow. Two copies of the same yellow, once brightness is normalized, remain yellow. Idempotence is a property of $\uplus$ alone: $A \uplus A = A$ always, whereas $A \boxplus A$ grows strength to $r' = \min(1, 2r)$ at unchanged hue, and returns the same yellow only after that normalization ([Axiom 5](#)). Two notes yield a chord, not a larger integer; a tone plus its opposite phase yields silence. In each case the word “add” names a different rule, and therefore a different kind of unit.

Thesis (an arithmetic of units of truth). *Mathematics need not be performed on numbers alone. Wherever a domain supplies a unit of truth — the smallest thing it treats as one — and a small set of rules by which such units combine, obeyed consistently within that domain, the unit and its rules together form an arithmetic of that domain. Counting is one instance: the unit is a token and the rule is $+$. Such rules are seldom invented. They are discovered, as the outcomes a domain returns every time the same conditions are repeated and it is left to settle toward stability.*

That is what an arithmetic *is*. What it is *for* is a result reached by a method that is consistent, reliable, and repeatable under the same conditions. Counting is one such method; it is not the only one. A hue, a chord, or a well-formed word can be such a result on another unit. This paper supplies the unit and the verbs. It does not, by itself, issue a verdict: a bowl may be mixed, walked, or assembled and brought to rest. Saying what that state means is a later reading, not a sixth verb here.

The technique is not new. Wherever a domain already treats something as one unit and combines those units under a small repeatable rule, people have been doing an arithmetic of that domain. Pāṇini’s grammar is the sharp linguistic case: roots and operations generate Sanskrit word-forms; commentators have called that engine mathematical, but only loosely. The kinship is the habit — unit plus rule, same conditions, same well-formed whole — not an identity with $\uplus$ or T . Placing a chair at a table is the weakest picture and is kept only as a reminder that the same two units, combined the same way, can settle to the same usable whole.

Examples are everywhere once looked for. Light combines additively and pigment subtractively: one domain, color, carries two arithmetics, because it is the rule set and not the domain that defines one. Notes combine into chords; a plate driven at the same frequency settles into the same nodal figure (§2.9); water freezing under the same conditions returns the same six-fold symmetry. The oldest instances are older than any notation for them. Atoms combine only in the proportions their valences allow, with mass and charge conserved across the reaction. Conserved quantities in particle physics, electric charge among them, add across every interaction, and some combinations are forbidden outright. In biology, an engulfed bacterium and its host became one cell, and the mitochondrion still carries a genome of its own: two units joined into a stable whole without either being dissolved — the shape of assembly ([Axiom 6](#)), not of a blend.

These domains are cited as evidence that such arithmetics exist, not as domains this paper models; [Axiom 0](#) decides which sorts earn a geometric reading here, and

most physical domains do not. Ice is the plainest case: its six-fold growth is a genuine arithmetic of its own domain, recorded by Axiom 0 as the snowflake’s native law, and it is exactly the kind refused a map onto this sphere. This paper states its own rules as axioms rather than discovering them. The second half of the thesis appears in it only in miniature: in the enclosure of Section 9, units left under one fixed law settle toward equilibrium and are read at a stated checkpoint.

When the object of interest is a structured whole rather than a count, a hue, or a pitch, the natural unit is a *pattern*. This paper represents one pattern by one typed point on a unit sphere $\mathcal{S} \subset \mathbb{R}^3$, not by a field painted over the globe:

- $r \in [0, 1]$ is the strength of the unit, capped so that a state remains a single unit of truth — an activation, not a global probability density;
- $\hat{n} \in S^2$, or equivalently (θ, ϕ) , is the direction of the pattern on the sphere;
- $\chi \in [0, 2\pi)$ is an internal phase-color, so two identities may occupy the same direction without overwriting one another;
- τ is the sort (**word**, **num**, **color**, ...): a legality tag, not a fifth geometric angle.

Pattern Arithmetic is the algebra of these units. Bundle does not double identical states ($A \uplus A = A$). A Bundled unit may be removed ($\{A\} \setminus A = \emptyset$). Rotations of the sphere preserve the unit. Opposite-phase MixMerge is a later reading; it is not $\uplus$.

A single numerical picture is enough to see the difference from counting. Take two unit states at the same place, one with phase 0° (red) and one with phase 120° (green). Counting would return 2. Phase-aware combination returns a single unit whose blended phase is 60° (yellow). That is this paper’s color inclusion — phasor MixMerge on a toy chart, not CIE or linear RGB ([Axiom 0](#), [Axiom 5](#)). The same two strengths with opposite phase (0° and 180°) return the null state, black. Two units of truth may become a new identity, a stronger mark of the same identity, or none.

The working order of the paper is not that MixMerge alone. One pipeline is

$$\text{Bundle (bowl)} \longrightarrow T(\text{path}) \longrightarrow \text{Result}.$$

The other is assemble ($\bowtie$) then rest of seats. Both use the same unit. Mixing is a later reading of a bowl, not a third spine. Collecting red and green is not yet yellow, and not yet a poem. Yellow is what you get if the later reading is MixMerge ($\boxplus$). Red and green may equally well walk a T_{poem} path; the result need not be yellow — though a given poem may still arrive there. The later act, not the shells, chooses the result. A verse is another path. Each word of “Twinkle, twinkle, little star...” is a unit state on the same sphere. Bundle yields the bag of distinct words (**twinkle** appears twice in the line and only once in the bowl). The poem is the path that walks those states in order, including the repeated **twinkle**. The bag is not the song. Numericity is not stored by $\uplus$. Two identical units Bundle to one address; that is the point of $A \uplus A = A$. How many times the address is *visited* lives in the path (the tape says **twinkle** twice). How many as a *count* lives in **num** mode, where $1 + 1 + 1 + 1 = 4$ and four tokens of 1 are not Bundled. Strength r is a poor substitute for either: a second **twinkle** cannot honestly raise r past 1, and “four ones” is schoolbook addition, not a fuller unit of truth. The same bowl may be read more than one way. Loading is not meaning.

The sphere is the stage. It is not the deepest claim. Strip away geometry, phase, rotations, and the language examples, and one proposition remains: different kinds of

information require different notions of combination, and an operation should be read by what informational property it *preserves* — and what it is allowed to throw away.

$$1 + 1 = 2, \quad \text{red} \boxplus \text{green} = \text{yellow}, \quad A \uplus A = A, \quad T_a(T_b(x)) \neq T_b(T_a(x)).$$

These are not alternative tricks for the same plus sign. They answer different questions:

Operation	Question
$+$	How many?
$\uplus$ / $\setminus$	What is present? / what is removed?
$\boxplus$ (MixMerge)	What composite state results?
T	What happens, and in what order?
offset	What walk is reused?
$\bowtie$	How do two units relate while both still sit?

Division as deconvolution or signal unmixing is not among these verbs; shown ill-posed for this algebra, not merely unbuilt (Section 7.2). Inverse rotation of T is well-defined. A MixMerge inverse is not.

A pattern can be treated as one unit. Identity and multiplicity come apart: $\uplus$ keeps kind and discards count; the path keeps visits; `num` keeps quantity. Composition and transformation are different verbs. Sorts may share a stage without sharing an operator. The geometry that follows is only a language in which those verbs can be written. It does not mint the relations; a lexicon or a learning rule must still seat Sun and Day. Section 1.4 keeps that claim narrow.

1.1 The Limitations of Discretized Matrix Grids

A discrete matrix grid is the wrong box for such a unit. Traditional matrix representations of continuous state spaces suffer from two severe structural bottlenecks:

- **Memory Scaling Bottlenecks:** A voxel cube of resolution N_{vox} stores $O(N_{\text{vox}}^3)$ cells, growing steeply with resolution.
- **Signal Identity Destruction:** Mapping multiple overlapping variables to identical spatial coordinates often results in the destructive overwriting of underlying signal identities.

1.2 The Phase-Colored Unit Sphere Framework

To resolve these grid constraints, the sphere is offered as one bounded object that can carry overlapping identities by direction and phase, without enlarging a lattice.

- **Spatial Orientation** $\hat{n} \in S^2$, or (θ, ϕ) : Defines continuous directional coordinates across the spherical surface.
- **Magnitude** $(0 \leq r \leq 1)$: Strength of that one unit, capped so the arrow remains a single unit of truth. An activation, not a density for the whole sphere.
- **Internal Color Phase** $(\chi \in [0, 2\pi))$: Serves as an independent phase degree of freedom.
- **Sort** τ : A legality tag (`word`, `num`, `color`, $\dots$). It does not move the arrow.

1.3 Core Contributions & High-Level Intuition

The principal contributions of Pattern Arithmetic are:

- **A typed unit:** one atom $x = (r, \hat{n}, \chi, \tau)$ on a bounded sphere, not a voxel cell and not a global probability field.
- **Admissibility of sorts (Axiom 0):** a sort earns a geometric χ - or T -reading only where a stated map preserves its native combination; most domains do not.
- **Verbs on that unit:** $\uplus$ collects identities ($A \uplus A = A$); $\boxplus$ reads a composite; T walks an ordered path on one register occupant; offset copies a relation; $\bowtie$ assembles two units into a relation without fusing them. Schoolbook $+$ is reserved for **num**.

Later sections distinguish a numeric encoding from a lexical encoding on the same sphere, so that a report may carry words and numbers together without forcing them through $\uplus$.

1.4 Scope of the claim

The actual claim is narrow. Counting ($1 + 1 = 2$) is not the only valid arithmetic for combining units of information. A combination rule should be chosen for what the unit *is* — a count, a color, a direction, a word — not applied by default because it is the arithmetic everyone already knows. The sphere, phase, and the operators $\uplus$, $\boxplus$, T , offset, and $\bowtie$ are one proposed way to make that claim precise and checkable. They are not themselves the claim, and they are not offered as a general theory of language, physics, biology, or computation. The paper does not attempt to compute the governing equations of those systems. It attempts to arrive at higher-order realities by reading many co-present units of truth — seated units, whether pinned by a lexicon or by a learned map — through a path, or by letting external seats rest under assembly. That is a computation of *presence, phase, order, and rest*, not a substitute for a field equation. The enclosure’s energy is a bookkeeping rule for seats $\vec{P}$; it is not a model of gravity or of matter. Storms, forests, a table and chair, and a pathogen cancellation later in the paper (Section 8) are costumes for the $\uplus$ -versus- T split: presence versus path. Related-work paragraphs (Section 10) draw boundaries against existing formalisms. None of that is a demonstrated model of weather, ecology, or epidemiology, or a replacement for quantum mechanics, embeddings, or transformers. A later widening of a costume into a domain claim would need its own evidence. It is not an extension of the illustration. Section 2.1 states this limit formally: a sort earns a geometric reading only where a structure-preserving map into MixMerge or T exists at a stated fidelity, and most physical domains do not offer one.

Every law in this paper is deterministic: the same starting units give the same result, step for step. Where a claim rests on a starting configuration, the configuration itself is printed, as in the worked examples of Sections 3 and 8 and the enclosure of §9.4. How a reader generates configurations of their own — by hand, or by any library under any seed — is left open; the printed one is the shared base. Printing it makes a result replayable. It does not make the result right.

The next section turns these definitions into formal mathematical axioms.

2 Axiomatic Foundations Of Pattern Arithmetic

2.1 Axiom 0: Admissibility of Sorts

The axioms that follow govern how a state $x = (r, \hat{n}, \chi, \tau)$ behaves once it sits on $\mathcal{S}$. They say nothing about which sorts τ have earned the right to sit there with a geometric reading attached. That prior question is this axiom, and it is separate from Table 2 in Section 5.4: the table says which *operators* two already-admitted sorts may use on each other; this axiom says which sorts may claim χ and T mean something about their own domain in the first place.

The test. A sort τ has a domain D_τ and, in the world, its own native way of combining two members of that domain, written $a \circ_\tau b$ for $a, b \in D_\tau$. Call τ *geometrically admissible* if there is a map

$$\varphi_\tau : D_\tau \longrightarrow \mathcal{S}, \quad \varphi_\tau(a \circ_\tau b) \approx \boxplus(\varphi_\tau(a), \varphi_\tau(b)) \quad \text{or} \quad T_{\varphi_\tau(a)}(\varphi_\tau(b)), \quad (1)$$

holding at a stated fidelity (exact, or approximate with the approximation named). Equation (1) is the Tier 1 test. Tier 0 admission, below, does not require it. If no such φ_τ exists, τ is *geometrically inadmissible* for Tier 1: it may still occupy a point on $\mathcal{S}$, but only under the inert treatment described below, never with a claim that χ or T tracks anything the domain itself does.

Three tiers.

- **Tier 0 — default, unstated.** `word` and `num` need no admissibility argument because neither makes a physical claim through the geometry. `num` has an exact, invertible *chart* (Section 5.2): decode, add in $\mathbb{Z}$, encode. That chart is not a MixMerge or T homomorphism — MixMerge on it is the clock trap of that same section, which is why MixMerge on `(num, num)` is type-refused. `word` makes no claim at all — T on a word is bookkeeping (order, presence), never a statement about linguistic physics. MixMerge applied to two word-states is inert geometry, not a χ -reading of language. Both are silently admitted at this inert level; nothing further needs to be said to use them.
- **Tier 1 — explicit inclusion.** A sort earns a genuine χ -reading only by stating its φ_τ and the fidelity it holds at. Musical pitch class is close to exact: octave equivalence makes pitch native to a circle, so χ for a note is close to literal, not borrowed. Color is close but not unique: additive light mixing is linear in some models and only approximately phasor-like in others, so an inclusion of `color` must say *which* reading it means — MixMerge under $\boxplus$ and MixMerge under linear RGB blend are different, non-interchangeable inclusions, and a paper that admits color must pick one and say so. This paper picks phasor MixMerge on a toy $0^\circ/120^\circ/60^\circ$ chart (Axiom 5). It does not pick CIE, and it does not pick linear RGB. A second Tier-1 sort, `chladni`, is worked in full in Section 2.9, where Equation (1) is met exactly within a mode and shown unavailable across modes.
- **Inadmissible — no φ_τ exists.** Some domains have a native combination law with no structure-preserving map to $(\boxplus, \boxtimes, T, \text{offset}, \boxtimes)$ at any useful fidelity. A snowflake’s native symmetry is the discrete point group D_{6h} , and two ice structures

meeting is a diffusion-limited accretion process governed by temperature and humidity gradients — not a continuous rotation in $\text{SO}(3)$. There is no honest $\varphi_{\text{snowflake}}$ satisfying Equation (1). Such a sort is not barred from the register outright: it may still be loaded as a point, but only demoted to the Tier-0 treatment already given to **word** — T used for order or presence only, no χ -reading granted, and no claim that a rotation of its point corresponds to anything the domain does. What is barred is Tier-1 status: an inadmissible sort may never be handed a stated φ_τ , because none would be true.

Sort	Native law $\circ_\tau$	MixMerge/ T map?	Tier
num	schoolbook +	no (chart only)	0
word	none claimed	n/a	0
note-pitch	pitch-class sum	near-exact	1
color	additive light	phasor MixMerge (here)	1
chladni	modal superposition	exact, within a mode	1 [§]
snowflake	D_{6h} growth	no	T1 no; T0 yes

Table 1: Admissibility is a property of the sort, decided before Table 2’s operator legality is even asked. Equation (1) is the Tier 1 test; Tier 0 admission does not use it. A sort absent from this table is undecided, not admitted; treat it as inadmissible until a φ_τ is stated. [§]**chladni** is admitted at Tier 1 with a stated restriction: exact for superposition within one mode, undefined across modes (Section 2.9).

This axiom is also what gives Section 1.4 teeth. “Not a general theory of physics, biology, or computation” is not only a promise of restraint; it is a consequence of Equation (1) failing for most physical domains a reader might otherwise expect the sphere to cover.

2.2 Process principle: Bundle, path, result

The A_i below are already seated units, each loaded from an observed sequence by E_L (Equation (42)); Bundle is the act that discards the order and count of that sequence, not the act that produces the units themselves. With that understood, Pattern Arithmetic proceeds in three named acts, kept separate so that a bowl may grow before it is read:

$$\underbrace{A_1 \uplus A_2 \uplus \cdots \uplus A_n}_{R = \text{the bowl}} \longrightarrow T_{w_k} \circ \cdots \circ T_{w_1}(x) \longrightarrow \text{Result.} \quad (2)$$

Here x is the register occupant; the w_i are actors, taken from the bowl or from the tape. T does not act on the set R as a set. $\uplus$ *collects*: identical tokens do not double ($A \uplus A = A$), and opposite marks may sit together until a later reading. T is an ordered path, not a MixMerge. The result is whatever that path (or another named reading) makes of the bowl — a poem, a report, a staged endpoint.

Paint is MixMerge ($\boxplus$), a reading of peer states at a common direction, not a T :

$$\text{red} \boxplus \text{green} = \text{yellow}.$$

The same pair may take T_{poem} instead. That result is a verse about the two colors, not a hue — unless the poem itself walks back to yellow. Shells do not own a pathway. A verse is another path. One bowl, several ways to read it. One may collect and read in a single gesture; the algebra does not force the merge. This principle is an axiom of *order of acts*, not a new formula for a 4-tuple.

Levels are closed. A result may be seated again. Equation (2) returns a state, and a state may take a seat in a further bowl, so the three acts repeat at a higher level: words Bundle and are walked into a line; the line is seated; lines Bundle and are walked into a verse. Two properties hold at every level and are worth stating once. First, a unit carries no representation of what it constitutes. A seated word does not know the poem, and T leaves τ_x invariant (Section 3.3), so being acted upon does not change what a unit is. Second, a bowl carries no representation of the reading that will be applied to it; that is what *one bowl, several ways to read it* means above.

The consequence is a design rule rather than a theorem. Whatever a lower unit would have to know about a higher one belongs at the higher level and not in the 4-tuple: which units were taken as co-present, over what interval, within what region, and under what question. Those are constituting facts of a level, not coordinates of an atom, and keeping them out is why the atom stays at four coordinates and one sort tag while the levels above it may be as elaborate as a problem requires. This paper works a single level and states the mechanism for the next; it does not carry a multi-level example.

2.3 Axiom 1: The Normalized Unit Truth State ($r \leq 1$)

A pattern is one typed point, not a field painted over the globe. The working atom is

$$x = (r, \hat{n}, \chi, \tau), \quad \begin{array}{ll} r \in [0, 1], & \hat{n} \in S^2, \\ \chi \in [0, 2\pi), & \tau \in \{\text{word, num, color, note-pitch, chladni, \dots}\}. \end{array} \quad (3)$$

Equivalently, $\hat{n}$ may be written in polar angles (θ, ϕ) , so the geometric part of the atom is the 4-tuple (r, θ, ϕ, χ) used by the operators below. The sort τ does not move the arrow; it says which buttons may be pressed (Section 5). A collection (bowl) is a separate finite set $B = \{x_1, \dots, x_n\}$ of such atoms. The sphere holds one current occupant of the register; the bowl is who is present.

The teaching atom is this 4-tuple: one arrow, $N = 1$. Axiom 7 allows an atom to hold N such arrows under one closure, each still capped at $r_i \leq 1$, plus one uncapped derived sum that is not itself a unit. That Collapse does not replace Equation (3). Every informational arrow in the Collapse is an instance of this axiom. Operators that need a single pair $(r, \hat{n})$ read the derived sum as a weight or a direction; as the actor of T in three dimensions, only after the cast of Axiom 7, and beyond three dimensions not from the derived sum at all, but from $N_{\text{dim}} - 2$ of the Collapse's own orthonormal arrows directly (§3.4). They do not mint a second kind of unit.

The cap $r \leq 1$ applies to *that* arrow. It is a normalized strength / activation coordinate, not a global probability density and not a demand that all arrows on $\mathcal{S}$ integrate to one. Two states may occupy different directions at once; each still carries its own strength $r \leq 1$. The sphere is the stage. The unit is the vector. $r = 1$ is a full-strength mark. Nothing in this axiom sums the globe into a single global budget. A field of many atoms, if wanted later, is a different object (for example a superposition of Dirac masses at those points), not the atom itself.

How the bound is imposed is a choice. This paper bounds by clipping: MixMerge returns $r' = \min(1, |z'|)$, so $r = 1$ is attained whenever a phasor sum reaches unit length. Clipping is not injective. Every sum of length one or more returns the same r' , and agreement past that point is not recorded (Section 2.7). An order-preserving bijection

from $[0, \infty)$ onto $[0, 1)$ — $|z| \mapsto |z|/(1 + |z|)$, say — would bound without discarding, since distinct sums would keep distinct strengths, and it would leave $z = 0 \mapsto 0$ and the black of [Axiom 2](#) untouched.

The clip is kept here, and the trade is stated rather than argued. Under such a bijection $r = 1$ becomes a limit and not a value: the saturated pole of [Axiom 3](#) would be approached and never occupied, and $\alpha = \pi r_w$ would never reach the half-turn of [Section 3.3](#). Two attained poles are worth more to this paper than lossless accumulation, because no construction here sums a large bowl. A setting that does — a verdict accumulated from many units, say — would weigh the same trade the other way, and should revisit this line rather than inherit it.

2.4 Axiom 2: The Absolute Null State ($r = 0$, Black)

The lower bound of a single unit is vanishing strength:

$$r = 0 \implies x = \mathbf{0} \quad (\text{Black}). \quad (4)$$

Direction and phase are then unused: there is no arrow to point. Black is an empty seat, not a field that is zero at every angle.

2.5 Axiom 3: Saturation (White)

White is the upper pole of the same bounded stage. It is not an infinite norm. Two readings are allowed; neither leaves the cap $r \leq 1$.

White as a colour. A single arrow may take a full-spectrum name (white light as one identity). That state is still one 4-tuple with $r \leq 1$.

White as a full house. The bowl may be saturated: many directions occupied, many hues present. Collecting further distinct shells fills more seats; it does not raise any arrow past $r = 1$. $\uplus$ does not grow strength. MixMerge may grow a single r until the cap. The stage looks white because the palette is on. The catalogue of names is large — on the order of a 24-bit palette, ~ 16 million colours per direction — and may be called practically endless. No unit is $r = \infty$.

$$\text{White} = \text{saturation of the bowl, or a full-spectrum name.} \quad (5)$$

Black is empty. White is as full as this construction allows. The menu is huge. The arrows stay finite. Density in this paper means many identities can share the stage without overwrite — many directions, and many phases per direction — each still capped at $r \leq 1$. It does not mean infinite information, and it is not a storage theorem against a voxel grid.

One atom carries a modest, bounded amount before it runs out of room. Two routes exist for more, and they are not the same route. The first is to seat a second atom: $\uplus$ gathers it into the same bowl ([Section 2.6](#)), which is how this construction represents anything beyond what one vector alone can hold, and is the route already in use throughout this paper. The second is to raise the ambient dimension, $S^2 \rightarrow S^{N_{\text{dim}}-1}$ ([Section 11.1](#)), giving a single atom more room on the sphere it already occupies. The second route is conditional on the first: it relieves crowding among many co-present atoms, and a bowl too small to crowd gains nothing by it — two atoms do not need a millionfold larger stage to stay apart.

2.6 Axiom 4: Bundle and Remove ($\uplus$, $\setminus$)

$\uplus$ *collects* a unit into the bowl. It does not MixMerge colours and it does not turn an arrow. A bowl is a finite subset of $\mathcal{S}$, and $\uplus$ is union on such subsets:

$$\uplus : 2^{\mathcal{S}} \times 2^{\mathcal{S}} \longrightarrow 2^{\mathcal{S}}, \quad B_1 \uplus B_2 = B_1 \cup B_2.$$

A bare state A appearing in an $\uplus$ -expression denotes its singleton bowl $\{A\}$. With that convention, identical names occupy one seat:

$$A \uplus A = A, \quad \text{that is,} \quad \{A\} \cup \{A\} = \{A\}, \quad (6)$$

while distinct names take separate seats, $A \uplus B = \{A, B\}$, so opposite marks may sit together: $\{+1\} \uplus \{-1\} = \{+1, -1\}$. Idempotence, commutativity and associativity of $\uplus$ are those of union, and none of the three needs separate postulation. Bundle is sort-blind: any two seated units may share a bowl, whatever their sorts, since a bowl records presence and presence makes no claim about a domain (Table 2). Sort discipline governs what may later be *done* with the bowl, not who may sit in it. The bowl is a set of states, not a blended 4-tuple.

$\setminus$ *removes* a unit, $\setminus : 2^{\mathcal{S}} \times \mathcal{S} \rightarrow 2^{\mathcal{S}}$. Removing the last occupant leaves the empty bowl:

$$\{A\} \setminus A = \emptyset. \quad (7)$$

The empty bowl $\emptyset$ and the null state $\mathbf{0}$ of Axiom 2 are distinct objects: $\emptyset$ is a bowl with no occupant, $\mathbf{0}$ is an occupant with $r = 0$. Only $\mathbf{0}$ is a state, which is why $T_{\mathbf{0}}$ is defined (Section 3.3) and $T_{\emptyset}$ is not. $\setminus$ undoes a Bundle. It is not anti-phase interference. Idempotence then gives the counting-trap

$$A \uplus A \uplus A \uplus A \setminus A = \emptyset \quad (\text{the empty bowl}). \quad (8)$$

Four gathers of the same name are still one seat; one removal empties it. Schoolbook $1 + 1 + 1 + 1 - 1 = 3$ lives only in **num** mode.

2.7 Axiom 5: MixMerge as a reading ($\boxplus$)

A blended 4-tuple appears only when MixMerge *reads* the bowl. Write $\text{MixMerge} = \boxplus$. MixMerge is a *phasor* operator: it combines χ among states that already share a direction $\hat{n}$. Direction belongs to T and to the rotations of §2.10; MixMerge never moves an arrow onto a new axis. MixMerge of states with differing $\hat{n}$ is left open (Section 11.1). The output is $(r', \theta', \phi', \chi')$ with $r' \leq 1$, or black. T_w is a one-argument rotation and phase-shift of an occupant, parameterized by an actor w . MixMerge is a combination of peer states. It is not a member of the group T_w lives in, and it does not compose with T_w the way two ordinary T 's do. The two give different answers on the same pair: for unit red at $\chi = 0^\circ$ and unit green at $\chi = 120^\circ$ sharing a direction, $T_{\text{green}}(\text{red})$ shifts phase additively to $\chi' = 120^\circ$, while $\text{red} \boxplus \text{green}$ takes the phasor resultant, $\chi' = 60^\circ$. No choice of actor makes T return 60° here, because T adds phase where MixMerge sums phasors. Three further properties separate them: T_w is asymmetric in actor and patient while $\boxplus$ is commutative; T_w leaves r invariant while $\boxplus$ may grow it to the cap (Axiom 3); and every T_w has an inverse w^* while $\boxplus$ has none.

Which sorts may read a MixMerge result as a claim about their own domain is settled by Axiom 0, not here: Tier 1 for **color** and **note-pitch**, inert geometry for **word**, type-refused for **num**.

Two states at a common direction. Treat strength and phase as a phasor $z = re^{i\chi}$. A domain reading requires A and B to share a sort (§2.1 above), so the sort carries through unchanged:

$$z' = r_A e^{i\chi_A} + r_B e^{i\chi_B}, \quad r' = \min(1, |z'|), \quad \chi' = \arg(z'), \quad \tau' = \tau_A = \tau_B, \quad (9)$$

and $(\theta', \phi') = (\theta_A, \phi_A)$. If $z' = 0$, the result is black. Equal unit paints at 0° and 120° give $z' = e^{i\pi/3}$, so

$$(1\angle 0^\circ) \boxplus (1\angle 120^\circ) = (1, \theta, \phi, 60^\circ, \text{color}) \quad (\text{yellow}).$$

This paper's color inclusion (Axiom 0, Tier 1) is that phasor MixMerge on a toy $0^\circ/120^\circ/60^\circ$ chart. It is not CIE, and it is not linear RGB. The unequal pair $r = 0.6$ at 0° and $r = 0.8$ at 120° gives $z' = 0.20 + 0.6928i$, hence

$$r' \approx 0.721, \quad \chi' \approx 73.9^\circ.$$

Equal anti-phase marks, same place, give $z' = 0$ (black). A state mixed with itself at the same phase is not idempotent: $A \boxplus A$ gives $z' = 2re^{i\chi}$, so $r' = \min(1, 2r)$. Strength grows until the cap clips it. That is distinct from $A \uplus A = A$.

Reading a whole bowl. MixMerge over a bowl B whose members share a direction $\hat{n}$ and, for a domain reading, a single sort τ is defined in one step, not as a fold:

$$\boxplus(B) = (\min(1, |z_B|), \hat{n}, \arg z_B, \tau), \quad z_B = \sum_{x \in B} r_x e^{i\chi_x}. \quad (10)$$

The binary $\boxplus$ of Equation (9) is the case $|B| = 2$. The one-step form is necessary, not stylistic, because the cap makes the binary operator non-associative: with $A = 1\angle 0^\circ$, $B = 1\angle 0^\circ$, $C = 1\angle 180^\circ$,

$$(A \boxplus B) \boxplus C = \text{black}, \quad A \boxplus (B \boxplus C) = 1\angle 0^\circ,$$

since $A \boxplus B$ is clipped from $2\angle 0^\circ$ down to $1\angle 0^\circ$, and the strength discarded by that clip is exactly what C would otherwise have cancelled. Equation (10) on the same three states sums first and caps once:

$$z_B = e^{i0} + e^{i0} + e^{i\pi} = 1, \quad \boxplus(\{A, B, C\}) = (1, \hat{n}, 0^\circ, \tau),$$

one reading, reached without choosing an order. $\boxplus$ is commutative, and Equation (10) is order-independent, so a bowl — which carries no order — has exactly one MixMerge reading. Iterated binary MixMerge does not, and is not what Equation (10) means. Non-associativity is the price of the cap of Axiom 1: without $\min(1, \cdot)$ the phasor sum would be ordinary complex addition and associativity would be free.

$\uplus$ does not run this rule. Two shells in a bowl have no resultant angle until MixMerge (or a T -path) is applied. A reading returns a state and leaves the bowl intact: $\boxplus(B)$ does not consume B , and a T -path draws actors from a bowl without removing them. Only $\backslash$ takes a member out. That is what makes one bowl readable more than once — the same B may yield a MixMerge result and, separately, a walked path. Bundle is lossless and is undone by $\backslash$; MixMerge is lossy, since it discards both order and any strength above the cap. That is why $\backslash$ inverts $\uplus$, w^* inverts T_w , and nothing inverts $\boxplus$.

Saturation and levels. A sum carries no bound of its own; the cap of [Axiom 1](#) supplies one. A bowl of many agreeing units therefore reaches $r' = 1$ quickly, and further agreement past that point is not recorded. This is worth reading as a remark about levels rather than a defect in the cap. A phasor sum is a reading of units that stand at one level together, and repeated samples of a single unit do not stand at one level; a bowl that saturates early is often a level that has been flattened ([Section 2.2](#)). Strength does not accumulate across levels: a reading returns a state whose r is legal by construction — $\min(1, |z_B|)$ for MixMerge, and the occupant’s own r for a path — and that state is seated afresh at the level above. Composition hides the question, since a path over many actors stays in $SO(3)$ and never saturates, so it is MixMerge rather than the walk that makes a flattened level visible. What the cap does not supply is a way to average, which needs the count that $\oplus$ discards ([Section 2.6](#)); reducing repeated samples to one seated unit is therefore work for the level below, not for $\boxplus$.

2.8 Process principle: UnMix

Named here for completeness of the verb inventory, not newly proven: UnMix ($\boxminus$), an inverse of MixMerge, does not exist.

$\boxminus$: not possible.

This is not an open question awaiting a construction. Given only a Collapsed read, the sources are not merely hard to recover; they are not determined by the result at all, verified directly in [Section 7.2](#): two genuinely different input pairs — the paper’s own red-and-green pair, and a second pair at 10° and 100° — reach the identical phasor sum to machine precision, and a wrong assumed composition reconstructs the observed result exactly, its residual indistinguishable from the correct pair’s own residual of zero. An algebra that states only forced, verified identities has no honest place for an operation whose wrong answer looks exactly like its right one. $\boxminus$ is recorded as a name for this already-settled fact, standing beside Collapse’s own UnCollapse so that the full inventory — Bundle, Remove, Collapse, UnCollapse, MixMerge, UnMix — has one entry per verb, proven or disproven, with nothing left unnamed.

2.9 A worked Tier-1 inclusion: chladni

Color is admitted on a declared toy chart. A second sort is admitted here on stronger ground, to show that [Equation \(1\)](#) can be met exactly rather than approximately, and by physics unrelated to either color or language.

Native law. A thin plate driven at one frequency responds as a superposition of its eigenmodes $\varphi_j(x, y)$, each carrying a complex amplitude $A_j = a_j e^{i\chi_j}$; the visible figure is the nodal set of $w = \sum_j A_j \varphi_j$. So $\circ_{\text{chladni}}$ is complex addition of mode amplitudes. Two drives on the same mode add as complex numbers.

Chart. One state is one mode. Let $\hat{n}$ carry the mode index under any injective map $(m, n) \mapsto (\theta, \phi)$, let χ be that mode’s phase, and let $r = a/a_{\max}$ be its amplitude normalized to a declared saturation $a_{\max}$.

The map, and where it is exact. Within one mode, $\hat{n}$ is fixed, and

$$\varphi_{\text{chladni}}(A \circ B) = \varphi_{\text{chladni}}(A) \boxplus \varphi_{\text{chladni}}(B) \quad (11)$$

holds identically while $|z| \leq 1$, because complex addition of amplitudes and phasor addition of (r, χ) are the same operation. Nothing is fitted. Note also that the chart's arbitrariness costs nothing here: same-mode combination never moves $\hat{n}$, so Equation (11) places no constraint on the mode-to-direction map. Two consequences are physical and immediate. Equal amplitudes in anti-phase give $z = 0$: the mode is driven to silence and its figure disappears, which is Equation (9) returning black. Two drives at 0° and 120° give $r = 1$ and $\chi' = 60^\circ$ — the same arithmetic as red $\boxplus$ green = yellow, on unrelated physics.

The restriction, which is physical. Across modes the map is not available, and not for want of effort. Mode space is not closed under addition: a sum of two distinct modes is not a mode, so it cannot be seated as a single state under any chart. Cross-mode coexistence is therefore $\uplus$ — a bowl of modes — and the resulting figure is a reading of that bowl which is not $\boxplus$. This is the open question of Section 11.1 met in a concrete domain, and it suggests the question may be closed negatively for this sort rather than left waiting.

Fidelity claimed, and three limits. Exact in phase and amplitude below saturation, within one mode; undefined across modes. Three limits are stated rather than argued away. The cap $r \leq 1$ is a normalization to a declared a_{max} , not a fact of linear plate theory, so Equation (11) clips above saturation, exactly as the color inclusion does. Direction is a label here and nothing more, so no meaning is claimed for T on a `chladni` state: rotating a mode index is not an operation the domain performs. And the figures most often reproduced are degenerate-mode combinations, which fall on the cross-mode side; what is admitted exactly is the simpler case, not the famous one.

2.10 Rotation preserves the unit

A rotation $\mathbf{R} \in \text{SO}(3)$ turns the direction $\hat{n}$ and leaves strength, phase, and sort alone:

$$\mathbf{R} \cdot (r, \hat{n}, \chi, \tau) = (r, \mathbf{R}\hat{n}, \chi, \tau), \quad \|\mathbf{R}\hat{n}\| = \|\hat{n}\| = 1. \quad (12)$$

The unit of truth is preserved because r is not scaled and $\hat{n}$ stays on S^2 . This is the geometric engine of T , not a second blend.

2.11 Axiom 6: Assembly ($\bowtie$)

Bundle lists who is present; MixMerge fuses peers into one state; T walks an ordered path. A fourth question remains: given two units, how strongly do they belong together, and where do they touch? Assembly, written $\bowtie$, answers it by adding a relation while leaving both units intact — it is neither $\uplus$ (which states presence and no relation) nor $\boxplus$ (which fuses and loses both identities):

$$A \bowtie B = (A, B, \kappa_{AB}, U_{AB}). \quad (13)$$

Both operands survive unchanged. What is new is the pair (κ_{AB}, U_{AB}) : a contact point and an interaction energy.

The interaction energy.

$$U_{AB} = -\frac{r_A r_B}{2} \cdot \frac{3 + \hat{n}_A \cdot \hat{n}_B}{(d_{AB}^2 + \varepsilon_{AB}^2)^{3/2}}, \quad \varepsilon_{AB} = r_A + r_B. \quad (14)$$

Three factors, each answering a different question. The product $r_A r_B$ asks how strong the two units are; a weak partner lowers the whole. The dot product $\hat{n}_A \cdot \hat{n}_B$ asks how well their directions agree, and is bounded: the bracket ranges over $[2, 4]$ only, so alignment can shift the result by at most a factor of two. The denominator asks how far apart they sit, which requires the enclosure of Section 9 and is the one term that reaches outside the bounded stage of Axiom 1. The softening length ε_{AB} is derived, not chosen: reading r as a radius, two units cannot approach closer than the sum of their own extents.

The contact point.

$$\kappa_{AB} = \hat{n}_A + \hat{n}_B, \quad (15)$$

the raw, unnormalized sum. Leaving it unnormalized is deliberate. The normalized bisector $\kappa_{AB}/|\kappa_{AB}|$ is undefined when $\hat{n}_B = -\hat{n}_A$, and that degeneracy drove a long search for a special case that turned out not to be needed: U_{AB} never required an axis at all, and κ_{AB} is a plain vector sum, defined at every angle. At exactly antipodal orientation it is $\mathbf{0}$ — a computed answer, not a failure, and the same absence of a favoured direction that Axiom 2 reports as black. Its magnitude is a free confidence signal, falling smoothly from 2 (identical directions) to 0 (opposed).

Assembly never repels. For two units of nonzero strength considered alone, $U_{AB} < 0$ always: there is no orientation, no pair of strengths, and no separation at which isolated assembly pushes rather than pulls. This is not an assumption but a property of Equation (14), whose bracket cannot fall below 2. The one boundary is a black unit (Axiom 2): at $r_A = 0$ the energy is exactly zero, no relation being formed with something absent, which is what black already means elsewhere. This is taken as the definition, $U_{AB} = 0$ whenever $r_A r_B = 0$, because the formula alone gives $0/0$ for two black units on one seat. Whether a structure of three or more units can hold members apart is a question for the enclosure, not for the pairwise verb.

Polarity. The difference $|r_A - r_B|$ distinguishes a balanced bond from a lopsided one, and so supplies a criterion for reading $A \bowtie B$ as mutual or as actor-and-patient in the manner of T . It is a label on the relation, not a difference in what the relation does: Equation (14) is symmetric under exchange of A and B , and $U_{AB} = U_{BA}$ exactly.

Where assembly sits among the verbs. Assembly is the first operator in this paper that requires anything outside the 4-tuple. Bundle, Remove, MixMerge, rotation, and offset all take units on the sphere and return units or states on the same sphere; none of them asks *where* a unit is, because nothing in $(r, \hat{n}, \chi, \tau)$ answers that. Assembly asks, and so is stated here as a verb but completed in Section 9, which supplies the seat $\vec{P}$ that d_{AB} measures between. Read without an enclosure, Equation (14) still gives a well-defined ordering of which pairs favour each other most strongly at equal separation.

2.12 Axiom 7: Informational vectors and derived per-atom metrics

The teaching atom of [Axiom 1](#) is one 4-tuple ($N = 1$). An atom may also be a *Collapse*: N informational vectors $v_i = (r_i, \hat{n}_i, \chi_i)$, each carrying its own phase-color $\chi_i \in [0, 2\pi)$ alongside $(r_i, \hat{n}_i)$ and sharing the atom's own sort τ — v_i together with that shared τ forming, per vector, the same 4-tuple [Axiom 1](#) states, still capped at $r_i \leq 1$. N is the atom's info number. This is not a second species of unit and not a field on the globe. It is several teaching atoms held under one closure — distinct phase-colors among them coexist without being forced through MixMerge, the Collapse's own instance of two identities occupying one direction without overwriting one another ([Axiom 1](#)). At most one further, *derived* vector may sit alongside them in the same closure unit:

$$V_{\text{derived}} = \sum_{i=1}^N r_i \hat{n}_i,$$

the raw, unnormalized sum of the informational vectors' geometric part, computed after and from them, never including itself; χ_i does not enter this sum, the same way τ does not participate in T 's rotation — available wherever a per-vector reading is named, silent in the seven metrics below, none of which is stated in terms of it. V_{derived} is uncapped and excluded from N : it needs no τ -justified admissibility of its own ([Axiom 0](#)), the same status the companion paper² already grants its own uncapped witness-sums. Its magnitude and direction read as an effective $(r, \hat{n})$ pair — $r = |V_{\text{derived}}|$, $\hat{n} = V_{\text{derived}}/r$ — standing in for the atom wherever a weight or a direction is required: Assembly's r_{AB} ([Axiom 6](#)), a camp body's alignment — neither involves a rotation or $\mathcal{S}$ -membership. It may act as the actor w of T (§3.2, whose domain is stated as $w \in \mathcal{S}$) only after an explicit cast back into $[0, 1]$ at the call site — the clip already used for MixMerge's own uncapped sum ([Axiom 5](#)), or, with no new machinery at all, atom coherence itself, $|V_{\text{derived}}|/S_{\text{atom}}$. Uncast, it is not a valid w : three vectors at $r = 0.5$, perfectly aligned, give $|V_{\text{derived}}| = 1.5$, outside $\mathcal{S}$ entirely, and $\alpha(1.5) = 270^\circ$ is not a rotation this axiom's T takes — not because $\text{SO}(3)$ forbids it, but because $w = 1.5$ was never T 's to accept. Until a cast is named at the call site, T takes teaching atoms only; every worked T path in this paper takes one, and a Collapse-as-actor example is not claimed here. For $N = 1$, V_{derived} coincides with the atom's own vector, already in $\mathcal{S}$, and nothing changes.

Seven scalar metrics follow from the same N informational vectors, each answering a different question, none used by this Axiom's own Assembly (§2.11) or by Motion (Section 9) unless named:

$$S_{\text{atom}} = \sum_{i=1}^N r_i, \quad \mathcal{V}_{\text{atom}} = \sum_{i=1}^N r_i^3, \quad \gamma_{\text{atom}} = \frac{\mathcal{V}_{\text{atom}}}{N}.$$

S_{atom} , atom strength, is the correct denominator for atom-level coherence, $|V_{\text{derived}}|/S_{\text{atom}}$, structurally identical to the companion paper's own mean resultant length one level down — reading 1.0 for perfectly aligned informational vectors and falling smoothly toward 0 as they scatter. $\mathcal{V}_{\text{atom}}$, atom volume, is strictly smaller than S_{atom} whenever any $r_i < 1$,

²*Pattern Arithmetic: Inference from Bowls and Paths*, which builds inference machinery — reference atoms, classification, this same τ -exempt status for uncapped witness-sums — on top of the atoms and operators defined here.

a forced consequence of the $r \leq 1$ cap rather than an assumption: cubing a number below 1 always shrinks it, exactly as treating r as a literal radius and r^3 as the volume it encloses would predict. γ_{atom} , mean vector volume, is the average per-vector share of that total. None of the three participates in Motion; atom volume, in particular, is deliberately excluded from it and exists here purely as a descriptive metric of the atom. Earlier drafts called $\mathcal{V}_{\text{atom}}$ “mass,” γ_{atom} “density,” and S_{atom} “atom energy” — all three retired as misleading: the first two because neither participates in momentum, kinetic energy, or any calculation the word “mass” would ordinarily license, and the third because a genuine notion of energy exists below and is not S_{atom} .

A fourth quantity is not derived from strength by squaring it, but stands on its own, the way variance is defined directly rather than as squared standard deviation:

$$\mathcal{E}_{\text{atom}} = \sum_{i=1}^N r_i^2. \quad (16)$$

$\mathcal{E}_{\text{atom}}$ is the combined squared magnitude the atom’s informational vectors would carry if every pair were mutually orthogonal — confirmed by direct construction, computing $\mathcal{E}_{\text{atom}}$ from a set of strengths matches $|V_{\text{derived}}|^2$ computed from those same strengths placed at mutually orthogonal directions, exactly. This is forced, not chosen: an exact identity places $\mathcal{E}_{\text{atom}}$ inside the derived vector’s own squared magnitude,

$$|V_{\text{derived}}|^2 = \mathcal{E}_{\text{atom}} + 2 \sum_{i < j} r_i r_j (\hat{n}_i \cdot \hat{n}_j),$$

verified to six decimal places. The structure matches the variance of a sum, $\text{Var}(\sum X_i) = \sum \text{Var}(X_i) + 2 \sum_{i < j} \text{Cov}(X_i, X_j)$, term for term: r_i^2 stands where $\text{Var}(X_i)$ stands, each vector’s own self-contained contribution; $r_i r_j (\hat{n}_i \cdot \hat{n}_j)$ stands where $\text{Cov}(X_i, X_j)$ stands, how two vectors relate. $\mathcal{E}_{\text{atom}}$ is the part that remains when every cross term vanishes — the atom’s intrinsic energy, prior to any question of how its own vectors align. This is distinct from the companion paper’s own bowl energy (a sum across many separate witnesses): bowl energy sums across atoms, $\mathcal{E}_{\text{atom}}$ sums squares within one.

A fifth metric answers a question the other four cannot. Atom coherence above is a ratio of magnitudes and so cannot fall below zero: it reads 0 both for informational vectors that merely fail to reinforce and for vectors that actively cancel. The strength-weighted mean cosine similarity across all internal pairs separates these:

$$\alpha_{\text{atom}} = \frac{\sum_{i < j} r_i r_j (\hat{n}_i \cdot \hat{n}_j)}{\sum_{i < j} r_i r_j}. \quad (17)$$

Each $\hat{n}_i \cdot \hat{n}_j$ is already a cosine similarity, since informational vectors are unit-directed; the denominator makes the aggregate comparable across atoms of differing N and differing strengths, which the unnormalized sum is not — two atoms with comparably tight internal arrangement but $N = 2$ and $N = 5$ show how little the raw sum alone says: $r = 0.85$ throughout, pairwise directions 5.7° apart, gives a raw sum of $+0.72$ at $N = 2$ (one pair) against $+7.14$ at $N = 5$ (a five-fold symmetric cone of the same half-angle, ten pairs) — a tenfold spread driven purely by pair count, not by any real difference in agreement, which α_{atom} correctly reports as 0.995 and 0.988, nearly identical. $\alpha_{\text{atom}} = +1$ exactly when all informational vectors align; it is negative when they genuinely oppose, reading $-1/3$ for two antipodal pairs on orthogonal axes (four unit vectors, six pairs, two at

cosine -1 and four at cosine 0). Its floor is not -1 in general: three or more vectors in three dimensions cannot be mutually opposed, so achievable disagreement tightens as N grows while achievable agreement does not — the unnormalized sum has maximum $\sum_{i<j} r_i r_j$, growing as N^2 , against a hard minimum of $-\frac{1}{2}\mathcal{E}_{\text{atom}}$, growing only as N , which follows directly from $|V_{\text{derived}}|^2 = \mathcal{E}_{\text{atom}} + 2\sum_{i<j} r_i r_j (\hat{n}_i \cdot \hat{n}_j) \geq 0$. Agreement compounds; disagreement saturates. α_{atom} is read alongside atom coherence, not in place of it: the first asks whether the vectors agree, the second how much of their strength survives into one resultant.

A sixth metric does for magnitude what α_{atom} does for direction. S_{atom} and $\mathcal{V}_{\text{atom}}$ both report a total and are silent on how that total is spread: two atoms can carry near-identical strength (2.200 against 2.150 in a checked case) while one holds it evenly across its vectors and the other concentrates most of it in two. Writing $s_i = r_i^3/\mathcal{V}_{\text{atom}}$ for the share of volume held by vector i :

$$\sigma_{\text{atom}} = \frac{N^2}{N-1} \cdot \text{Var}(s_1, \dots, s_N), \quad N \geq 2, \quad (18)$$

with $\sigma_{\text{atom}} = 0$ by convention at $N = 1$. $\sigma_{\text{atom}} = 0$ exactly when volume is spread evenly and 1 when a single vector holds all of it, at every N — verified at $N = 2, 3, 4, 6, 8$ for both extremes. The normalization is what earns this form over the simpler candidates: the largest vector’s share and the Herfindahl index were both tested and rejected, since each reads a perfectly even atom as $1/N$ and so slides from 0.5 to 0.125 as N grows from 2 to 8, confounding evenness with size. $\mathcal{V}_{\text{atom}}$ rather than S_{atom} supplies the shares deliberately: cubing amplifies concentration, so a single dominant vector at $r = 0.95$ against three at $r = 0.2$ registers a share of 0.973, where linear strength would report the same atom as unremarkable.

A seventh metric answers a question none of the first six can: not how strength or volume is shared among distinct vectors, but whether many vectors independently arrive at nearly the same magnitude. σ_{atom} is blind to this by construction — four hundred vectors at an identical $r = 0.65$ each receive their own tiny individual share, indistinguishable to Equation (18) from four hundred vectors scattered across the full range; checked directly, a genuinely clumped population registered as *more* even ($\sigma_{\text{atom}} = 0.00032$) than a uniform spread of the same size (0.00214). Fixing this needs a shell width w , a named, configurable parameter rather than a derived one, defaulting to $w = 0.1$ — ten shells spanning $[0, 1]$. Binning the N informational vectors by $\lfloor r_i/w \rfloor$ into K occupied shells with populations $n_1, \dots, n_K$ and $p_k = n_k/N$:

$$\sigma_{\text{shell}}(w) = \begin{cases} 1, & K = 1 \\ \frac{K^2}{K-1} \cdot \text{Var}(p_1, \dots, p_K), & K \geq 2. \end{cases} \quad (19)$$

The $K = 1$ case is stated separately because the general form has a removable singularity there: every vector sharing one shell is maximal concentration, not the zero the formula would otherwise return. $\sigma_{\text{shell}} = 0$ when occupied shells are evenly populated, at every K tested; $\sigma_{\text{shell}} = 1$ when every vector shares one shell. Checked at $w = 0.1$ and $w = 0.05$: well separated shells (as in 10@0.98, 3@0.3, 400@0.65, giving $\sigma_{\text{shell}} = 0.908$ at either width) are stable to the choice, while genuinely uniform populations are not — occupied shell count nearly doubled between the two widths in a checked uniform case, though the concentration score itself stayed negligible either way. w should be reported alongside σ_{shell} whenever it is used.

2.13 Process principle: Collapse

Bundle discards order and count on the way to a bowl. A second, separate process does the opposite: several already-seated units may combine into one atom, with every contribution kept. Call this Collapse ($\oplus$). Where Axiom 7 states what such an atom *is* — N informational vectors under one closure — this section states the operation that *produces* one from separately seated units, and settles a question that structure alone does not answer: what Collapse does with repetition.

$$A_1 \oplus A_2 \oplus \cdots \oplus A_n = A_{\text{collapse}},$$

where A_{collapse} is the atom of Axiom 7 holding $v_1, \dots, v_n$, one vector per contribution, each still the $(r_i, \hat{n}_i, \chi_i)$ of its own source, sharing the atom's one sort τ (Axiom 7's own present requirement; see the open item below). Collapse, like MixMerge, is a reading: it does not consume its sources. $A_1, \dots, A_n$ remain exactly as seated as they were, in a bowl or otherwise; Collapse produces an additional atom alongside them, not in place of them — the same principle already stated for MixMerge's own reading of a bowl, and the precondition that makes the identity-handle proposal below coherent at all: a pointer back to sources that had been consumed would point at nothing.

Collapse is not idempotent. $A \oplus A \neq A$ in general. Collapsing A with itself gives $N = 2$: two vectors, both reading A 's own $(r, \hat{n}, \chi)$, not one. Shorthand for this case, used below and throughout: $A^{(N)}$ for N identical copies, so $A \oplus A = A^{(2)}$. This is deliberate, and distinguishes Collapse from Bundle on the one point where structure alone leaves a choice open. Three considerations fix it.

First, Collapse's own stated purpose is that no contribution is lost. Were $A \oplus A = A$, the second contribution would be discarded exactly as Bundle already discards a repeated member ($A \uplus A = A$) — leaving nothing to distinguish the two verbs in this case, when the reason for having both is that one discards multiplicity and the other does not.

Second, idempotence here is inconsistent with UnCollapse under its own, otherwise solid base case (below): if $A \oplus A = A$, then undoing one contribution must return the other, $A \text{ UnCollapse } A = A$, contradicting $A \text{ UnCollapse } A = 0$ — which holds independently, matching $\{A\} \setminus A = \emptyset$ at the atom level. Both cannot stand unless $A = 0$.

Third, and most directly: a construction elsewhere in this paper already needed fixing for exactly this failure mode. Two independent witnesses that happen to agree exactly must remain distinguishable from one witness stated once, on pain of repeating the evidence-count blindness that a clip-versus-bijection choice was built to correct (companion paper, *Inference from Bowls and Paths*). Idempotent Collapse reintroduces that same blindness in a different verb.

Where N is kept. No new field. N is already Axiom 7's own count of the atom's informational vectors — “the atom's info number.” An atom that discards no contribution simply has as many v_i as contributions it received; N was never anything other than a property of that list.

Open: Collapse across differing sorts. Axiom 7 currently requires every vector in one atom to share the atom's single τ . Two atoms differing only in sort — equal r , equal $\hat{n}$, different τ — cannot yet be Collapsed under that requirement, and χ does not supply what is missing: χ 's stated role is distinguishing several identities already sharing one

direction (Axiom 1), not distinguishing sort. Admitting a per-vector τ_i would be the natural extension, paralleling χ_i 's own per-vector role, but it is not adopted here and is left open.

2.14 Process principle: UnCollapse

UnCollapse ($\ominus$) removes one contribution from a Collapsed atom, reducing its count by exactly one:

$$A_{\text{collapse}} \ominus A_k = A'_{\text{collapse}},$$

where A'_{collapse} holds the remaining $N - 1$ vectors. $N \rightarrow N - 1$ is the general rule, uniformly, at every N — UnCollapse is not a reset to $N = 1$ regardless of starting count; it removes exactly the one contribution named.

Two cases close the ladder at its ends. At $N = 1$, removing the atom's only contribution leaves black:

$$A \ominus A = 0,$$

matching $\{A\} \setminus A = \emptyset$ at the level of a single bowl member. At $N = 2$ down to $N = 1$, an atom holding one remaining vector *is* the teaching atom of that vector (Axiom 1's own $N = 1$ base case), so

$$A^{(2)} \ominus A = A^{(1)} = A,$$

correctly, because $2 - 1 = 1$ — not a general property of UnCollapse, but the one case where its arithmetic happens to land there. No further contradiction remains once Collapse's own idempotence is fixed above.

What UnCollapse requires to be well-defined. Removing " A_k " presumes a specific contribution is named, the same way $\{A\} \setminus A$ already presumes a specific member. Where every v_i in A_{collapse} differs in value, the named vector is unambiguous. Where two or more contributions share identical values — genuinely distinct sources that happen to agree exactly — $(r_i, \hat{n}_i, \chi_i)$ alone cannot distinguish which is meant. Resolving this needs an identity handle, id_i , carried alongside each vector, external to its geometric part the way a seat $\vec{P}$ is external to Axiom 1's own tuple: UnCollapse is value-exact without one, in every case where values alone suffice, and identity-exact with one, in the degenerate case where they do not. This mechanism is proposed, not adopted as an axiom here, and is left for the same later treatment as the seat of Section 9.

3 Transformation Algebra: From Identity to Expression

3.1 Motivation: Two Layers of Meaning

The operators established so far ($\uplus$, $\setminus$) answer the question *who is present?* Combining Sky, Sun, and Wind under $\uplus$ yields one order-independent bowl — a bag of seats, not a blended hue. Paint is MixMerge ($\boxplus$), a later reading. Analogies such as Sun-Day-Moon are *not* gathers; they are deferred to the offset operator of Section 4. Bundle $\uplus$ does not, and structurally cannot, distinguish a poem in which the sun rises through a wind-blown sky from a poem in which wind extinguishes a darkening sun — both share the identical bowl $R = \text{Sky} \uplus \text{Sun} \uplus \text{Wind}$, since $\uplus$ is commutative and idempotent by [Axiom 4](#).

The distinction between these two poems is not a difference in *ingredients*; it is a difference in *how one state acts upon another, and in what order*. This is a *relational* operation, and it requires a second operator, T , layered on top of $\uplus$: where $\uplus$ builds nouns, T builds sentences. A poem is not the bowl R itself, but a specific trajectory traced across $\mathcal{S}$ by successive transformations applied to a register occupant.

3.2 Formal Definition of the Transformation Operator

Rather than introduce T as an unrelated new primitive, it is defined using the same rotational machinery already given in §2.10: every state $w \in \mathcal{S}$ acts on another state $x \in \mathcal{S}$ by rotating x 's direction about w 's own axis, by an angle proportional to w 's strength, while additively shifting x 's phase by w 's phase:

$$T_w(x) = \left(r_x, \text{Rot}_{\hat{n}_w}(\alpha(r_w))\hat{n}_x, (\chi_x + \chi_w) \bmod 2\pi, \tau_x \right) \quad (20)$$

where $\hat{n}_w$ is the unit axis of w , $\text{Rot}_{\hat{n}}(\alpha) \in \text{SO}(3)$ is the rotation of §2.10 about that axis by angle α , and $\alpha(r_w) = \pi r_w$ is the simplest monotonic choice mapping strength to rotation angle (a full-strength unit, $r_w = 1$, rotates by the maximal 180°). Magnitude r_x and sort τ_x are left invariant: a transformation reorients and recolors a state, it does not strengthen or weaken it, and it does not recast it — those remain the roles of MixMerge (strength) and of an explicit cast (sort).

3.3 Derived Properties

The following properties follow from Equation (20) rather than being separately postulated.

Identity. If $w = \mathbf{0}$ (the Null State of Axiom 2), then $\alpha(0) = 0$ and the phase shift is 0, so $T_0 = \text{id}$.

Invertibility. The inverse of T_w is not “the same axis and the same $+\alpha$ again.” Same axis, same positive angle, repeats the turn. Inverse is $-\alpha$ on the same axis, or $+\alpha$ on the antipodal axis, together with $-\chi_w$. Define

$$w^* = (r_w, -\hat{n}_w, -\chi_w \bmod 2\pi, \tau_w), \quad (21)$$

where $-\hat{n}_w$ is the antipode ($\pi - \theta_w, \phi_w + \pi$). Then $T_{w^*}(T_w(x)) = x$ for all x . As rotations, only $\alpha = 0$ ($r_w = 0$) and $\alpha = 180^\circ$ ($r_w = 1$) equal their own reverse; the inverse state w^* is still required so that phase is cancelled too. This undoes a transformation on a state that persists. It is distinct from $\{A\} \setminus A = \emptyset$ (last shell removed, leaving the empty bowl) and from MixMerge of equal anti-phase marks (a reading into black). Neither of those is an inverse of T .

Associativity. Since each T_w is a function $\mathcal{S} \rightarrow \mathcal{S}$, composition of transformations is associative automatically, by the associativity of function composition; no additional axiom is required.

Non-commutativity. Rotations about distinct axes in $\text{SO}(3)$ generically do not commute: for a, b with distinct directions (θ, ϕ) ,

$$T_a(T_b(x)) \neq T_b(T_a(x)). \quad (22)$$

This is not an additional assumption; it is inherited directly from the well-established non-commutativity of $\text{SO}(3)$, already invoked by §2.10. It is this property, and only this property, that allows two poems built from the same bowl to diverge: order of application is meaningful even though the underlying identity is not.

3.4 Beyond three dimensions

Equation (20) needs only $\hat{n}_w$ to fix a rotation because $\text{SO}(3)$'s generators happen to form a 3-dimensional space, exactly matching $\mathbb{R}^3$ — not because one vector generally suffices to specify a rotation. It does not: a *simple* rotation, one that fixes as much as possible and turns a single remaining plane, fixes a subspace of dimension $N_{\text{dim}} - 2$ and rotates its 2-dimensional orthogonal complement. The ladder climbs one dimension at a time. A circle turns about a fixed centre; stacking that centre through the circular slices of a sphere gives the sphere's 1-dimensional axis; stacking again gives the 2-dimensional axis plane of the sphere in four dimensions, and again the 3-dimensional fixed space of the sphere in five:

Space	Sphere	Fixed part (the “axis”)	What turns
$\mathbb{R}^2$	S^1 , a circle	a point, dimension 0	one plane
$\mathbb{R}^3$	S^2	a line, dimension 1	one plane
$\mathbb{R}^4$	S^3	a plane, dimension 2	one plane
$\mathbb{R}^5$	S^4	a 3-D space, dimension 3	one plane
$\mathbb{R}^{N_{\text{dim}}}$	$S^{N_{\text{dim}}-1}$	dimension $N_{\text{dim}} - 2$	one plane

At every rung the fixed part grows by one dimension and the turning part stays a single plane. (A sphere is named by its own dimension, not by the space around it: the sphere in four dimensions is S^3 .) Three dimensions is not where axes stop needing more than a point — it is where they still fit in one vector.

For $\mathcal{S} = S^{N_{\text{dim}}-1} \subset \mathbb{R}^{N_{\text{dim}}}$, w supplies the fixed subspace F directly, as $N_{\text{dim}} - 2$ mutually orthogonal informational vectors from its own Collapse (Axiom 7) rather than as one $\hat{n}_w$:

$$T_w(x) = \left(r_x, \hat{n}_x^F + \text{Rot}_{\alpha(r_w)}(\hat{n}_x^P), (\chi_x + \chi_w) \bmod 2\pi, \tau_x \right) \quad (23)$$

where $\hat{n}_x^F$ is $\hat{n}_x$'s component in F , left untouched, and $\hat{n}_x^P$ its component in the remaining 2-dimensional plane $F^\perp$, rotated within that plane by $\alpha(r_w) = \pi r_w$ exactly as before. This requires $N_{\text{dim}} - 2$ mutually orthogonal arrows in the actor's Collapse, so $N \geq N_{\text{dim}} - 2$ is necessary but not by itself sufficient: an atom acting as w in more than three dimensions needs a Collapse rich enough to supply its own axis, not just a direction. At $N_{\text{dim}} = 3$, F is one vector and Equation (23) is Equation (20) exactly — checked directly, not merely argued: the same inputs give the same output to machine precision.

Stage	Orthonormal arrows required on the <i>actor</i>
S^2 (3-D)	1 — the teaching atom, unchanged
S^3 (4-D)	2
S^{11} (12-D)	10

A teaching atom, $N = 1$, is therefore a legal actor only at $N_{\text{dim}} = 3$; this is why three dimensions is the standard this paper keeps, not an arbitrary default — a single arrow can walk without a Collapse only there.

Checked beyond the coincidence of matching 3D: invertible by $-\alpha$ on the same w at every angle tested from 0 to 359° , not only small ones — the naive alternative, borrowing the missing axis data from x instead of w , is invertible only while α stays short of the angle between w and x , and silently fails past it, since the borrowed direction depends on which side of w the rotation has already passed. Sourcing F from w alone avoids this: the fixed subspace never depends on what is being acted on, the way $\text{SO}(3)$ ’s axis never did. Non-commutativity for distinct actors and angle-composition for repeated ones both continue to hold, checked directly at $N_{\text{dim}} = 4, 5, 8, 12, 20$, and 50 on a configuration that needs no seed. With e_i the coordinate axes, actor w_1 fixes $e_1, \dots, e_{N_{\text{dim}}-2}$ and so turns the plane of $e_{N_{\text{dim}}-1}, e_{N_{\text{dim}}}$; actor w_2 fixes $e_2, \dots, e_{N_{\text{dim}}-1}$ and so turns the plane of $e_1, e_{N_{\text{dim}}}$; each plane is oriented from its lower-indexed axis to its higher. The occupant is $\hat{n}_x = (1, 2, \dots, N_{\text{dim}})$ normalized. Invertibility is tested at $30^\circ, 45^\circ, 90^\circ, 135^\circ, 180^\circ, 200^\circ, 270^\circ$, and 359° ; the order gap is $|T_{w_2}T_{w_1}\hat{n}_x - T_{w_1}T_{w_2}\hat{n}_x|$ with both actors at 90° :

N_{dim}	largest inversion error	order gap
4	1.1×10^{-16}	0.683
5	1.1×10^{-16}	0.688
8	1.1×10^{-16}	0.649
12	5.6×10^{-17}	0.584
20	1.1×10^{-16}	0.489
50	5.6×10^{-17}	0.331

Inversion is exact to rounding at every angle tested, and the two orders disagree at every dimension tested. Both results stay on the unit sphere.

One choice is left open, and is a convention rather than a gap: which orthonormal basis of $F^\perp$ orients the rotation (which direction counts as $+\alpha$) is fixed by w ’s own $N_{\text{dim}}-2$ vectors, but the specific construction is not unique — no more than the right-hand rule is the only possible convention for $\text{SO}(3)$ ’s own cross product. Any consistent choice, stated once, suffices; §3.7 states and exercises one.

3.5 Poems as Transformation Trajectories

A poem π built from a register occupant x and an ordered sequence of transforming states $w_1, w_2, \dots, w_n$ (typically drawn from the bowl) is the trajectory

$$\pi = (x; w_1, \dots, w_n) : \quad x \longrightarrow T_{w_1}(x) \longrightarrow T_{w_2}T_{w_1}(x) \longrightarrow \dots \longrightarrow T_{w_n} \dots T_{w_1}(x). \quad (24)$$

Two poems may share an identical bowl (the same theme, the same “ingredients”) while tracing entirely different trajectories — this is the formal sense in which many poems arise from the same units of truth. Conversely, because rotation composition is many-to-one, two distinct trajectories may terminate at the same endpoint; a poem’s identity is therefore carried by its full path through $\mathcal{S}$, not by its final state alone.

Unlike $\oplus$, T does not Bundle a repeated actor: $T_w \circ T_w \neq T_w$ in general, since composing a rotation with itself composes the angle, $T_w \circ T_w = T_{w'}$ for some $w' \neq w$ rather than returning T_w . A path that visits the same actor n times is therefore n genuine steps,

not one step counted n times, and its length recovers that count without any coordinate carrying it. This is the sense in which the path, unlike the bowl, does not discard multiplicity (Section 2.6).

A step of the trajectory may carry more than the bare actor: $\pi = ((x; w_1, m_1), \dots, (w_n, m_n))$ for some label m_i — a discrete time bucket, say. This changes nothing about the atom: m_i is metadata on the walk, not a coordinate of any x or w_i , and Equation (24) runs exactly as before on the w_i alone. What such a label buys, and does not buy, follows directly from the discreteness point above: a handful of buckets behaves like a sort tag and does not block $A \uplus A = A$ where it applies, but a bucket is a fact about how a particular walk was assembled, not a property its occupant carries into the next one.

3.6 Worked Numerical Example: Sun, Wind, and Sky

Let Sky be the state acted upon, with direction $(\theta, \phi) = (45^\circ, 45^\circ)$, i.e. Cartesian axis $(0.500, 0.500, 0.707)$. Let Sun define an axis along $(\theta, \phi) = (90^\circ, 0^\circ)$, i.e. $(1, 0, 0)$, with strength $r_{\text{Sun}} = 1.0$, giving rotation angle $\alpha_{\text{Sun}} = 180^\circ$. Let Wind define an axis along $(\theta, \phi) = (90^\circ, 90^\circ)$, i.e. $(0, 1, 0)$, with strength $r_{\text{Wind}} = 0.7$, giving $\alpha_{\text{Wind}} = 126^\circ$.

Order 1 — Sun acts first, then Wind (the sun rising through a wind-blown sky):

$$T_{\text{Wind}}(T_{\text{Sun}}(\text{Sky})) \approx (-0.866, -0.500, 0.011) \quad (25)$$

Order 2 — Wind acts first, then Sun (wind extinguishing a darkening sun):

$$T_{\text{Sun}}(T_{\text{Wind}}(\text{Sky})) \approx (0.278, -0.500, 0.820) \quad (26)$$

The two endpoints are clearly distinct, confirming Equation (22): identical primitives, identical bowl, opposite structural readings. Notably, if $\chi_{\text{Sun}} = 0^\circ$ and $\chi_{\text{Wind}} = 90^\circ$, the accumulated phase shift is 270° in *both* orderings, since phase addition is commutative. The two poems therefore share the same blended color/mood even as their spatial — structural — meaning diverges. This asymmetry (shared affect, divergent structure) is a numerical check of the algebra, not a theory of verse.

3.7 Worked Numerical Example: A Four-Dimensional Actor

The construction of §3.4 is checked here by hand, not only by claim. Let $N_{\text{dim}} = 4$. An actor w supplies its required $N_{\text{dim}} - 2 = 2$ orthonormal arrows as the first two standard basis vectors, $e_1 = (1, 0, 0, 0)$ and $e_2 = (0, 1, 0, 0)$, fixing $F = \text{span}(e_1, e_2)$ and leaving the active plane $F^\perp = \text{span}(e_3, e_4)$. Take $r_w = 0.5$, so $\alpha(r_w) = 90^\circ$. Let the occupant be

$$\hat{n}_x = \left(\frac{1}{2}, \frac{1}{2}, \frac{1}{\sqrt{2}}, 0 \right),$$

a unit vector with a nonzero component in both F and $F^\perp$, so the example shows both halves of Equation (23) at once. Its F -component, $(\frac{1}{2}, \frac{1}{2}, 0, 0)$, is exactly what T_w leaves untouched; its $F^\perp$ -component, $(0, 0, \frac{1}{\sqrt{2}}, 0)$, is what turns by 90° within $\text{span}(e_3, e_4)$, landing on $(0, 0, 0, \frac{1}{\sqrt{2}})$. Summing:

$$T_w(\hat{n}_x) = \left(\frac{1}{2}, \frac{1}{2}, 0, \frac{1}{\sqrt{2}} \right). \quad (27)$$

Applying T_w again with $-\alpha$ recovers the $F^\perp$ -component exactly, $(0, 0, \frac{1}{\sqrt{2}}, 0)$, since F was never touched to begin with:

$$T_w^*(T_w(\hat{n}_x)) = \left(\frac{1}{2}, \frac{1}{2}, \frac{1}{\sqrt{2}}, 0\right) = \hat{n}_x, \quad (28)$$

confirming invertibility at a genuine right angle, not only for the small perturbations a reader might otherwise suspect were doing the work. The orientation convention of §3.4 is exercised here for the first time: $f_1 = e_3$, $f_2 = e_4$, in that order, is the specific choice this example makes, and a different ordering would rotate the other way.

3.8 Scope

T is deliberately left unconstrained beyond Equation (20): any sequence of any transforming states is a formally valid trajectory in three dimensions, or in N_{dim} dimensions provided each actor's own Collapse is rich enough to name its F (§3.4); an underqualified actor is not a shorter path, it is not a legal w at all. Pattern Arithmetic in this form supplies the substrate on which expression is built — the combinatorial space of possible trajectories — and does not yet supply a grammar constraining which trajectories are coherent, well-formed, or meaningful to a human reader. Whether such constraints belong inside this framework (as further axioms on admissible sequences of T) or outside it (as a learned or externally specified grammar layered on top) is left open.

4 Analogical Arithmetic: The Offset Operator

Offset is another verb. It is not $\uplus$ (no Bundle), not MixMerge (no phasor blend), and not T (no spin of an actor on a path). It copies a *walk* from one seat to another. The landing is a new direction on the same globe. A dictionary may later pin a name on that landing. The walk does not mint the name.

4.1 Recipe: logarithm, carry, exponential

Write directions as unit vectors $\hat{n} \in S^2 \subset \mathbb{R}^3$. The walk from p to q is the tangent vector

$$\theta = \arccos(\hat{p} \cdot \hat{q}), \quad \log_{\hat{p}}(\hat{q}) = \frac{\theta}{\sin \theta} (\hat{q} - (\hat{p} \cdot \hat{q}) \hat{p}) \quad (\theta \neq 0), \quad (29)$$

with $\log_{\hat{p}}(\hat{p}) = \mathbf{0}$. Its length is the geodesic angle. The reverse map, from a tangent vector v at p , is

$$\exp_{\hat{p}}(v) = \cos \|v\| \hat{p} + \sin \|v\| \frac{v}{\|v\|} \quad (v \neq \mathbf{0}), \quad (30)$$

and $\exp_{\hat{p}}(\mathbf{0}) = \hat{p}$. Given three seats m, k, w (source, source-image, new start), set

$$v = \log_{\hat{n}_m}(\hat{n}_k), \quad \hat{n}_* = \exp_{\hat{n}_w}(\text{transport}_{m \rightarrow w}(v)). \quad (31)$$

The carry rotates the tangent vector by the geodesic angle about the axis of $m \rightarrow w$:

$$\theta_{mw} = \arccos(\hat{n}_m \cdot \hat{n}_w), \quad \hat{u} = \frac{\hat{n}_m \times \hat{n}_w}{\|\hat{n}_m \times \hat{n}_w\|} \quad (\hat{n}_m \neq \pm \hat{n}_w), \quad (32)$$

$$\text{transport}_{m \rightarrow w}(v) = \text{Rot}_{\hat{u}}(\theta_{mw}) v = v \cos \theta_{mw} + (\hat{u} \times v) \sin \theta_{mw} + \hat{u} (\hat{u} \cdot v) (1 - \cos \theta_{mw}). \quad (33)$$

The axis is undefined when $\hat{n}_m = \pm \hat{n}_w$. Coincident seats ($\hat{n}_m = \hat{n}_w$): the carry is v itself. Antipodal seats ($\hat{n}_m = -\hat{n}_w$): no unique short geodesic; that case is a separate convention and is not used below. For a first recipe, if m and w are not far, transport may be omitted and v used at w as written. That case is crude and marked as such. The full recipe is (33), so the heading stays a ground heading.

Strength and colour are not invented by the walk. This paper takes

$$r_* = r_w, \quad \chi_* = (\chi_w + \chi_k - \chi_m) \bmod 2\pi, \quad \tau_* = \tau_w. \quad (34)$$

The target keeps its r and its sort. The mood increment of the pair $m \rightarrow k$ is reused. Nothing here uses $\alpha = \pi r$ from T . With (31) and (34) together, offset returns a complete state, not a bare direction:

$$\text{offset} : \mathcal{S} \times \mathcal{S} \times \mathcal{S} \longrightarrow \mathcal{S}, \quad (m, k, w) \mapsto (r_w, \hat{n}_*, \chi_*, \tau_w).$$

It is therefore a verb on the same footing as $\uplus$, $\boxplus$ and T : a landing may be Bundled into a bowl, walked from, or used as the start seat of a further offset. The question of *whether anything is seated there* is separate, and is settled by a dictionary, not by the walk. $\tau_* = \tau_w$ because the question asked is what stands to w as k stands to m ; the answer is of w 's kind. A walk copied between sorts would need an explicit cast (Section 5), which offset does not perform. Two seated units keep two strengths. If Day has $r = 0.9$ and Night has $r = 0.8$, both sit; there is no $0.9 + 0.8 = 1.7$. That sum is ambient $\mathbb{R}^3$ arithmetic and leaves the ball. $\uplus$ gathers the two seats. A later path may turn a dial or clip a MixMerge. It does not get to smash two r 's into one illegal length as the price of seating them.

Beyond three dimensions. Equation (33) is written with a cross product, which exists only in three dimensions, but what it performs is the simple rotation of §3.4: it turns the plane spanned by $\hat{n}_m$ and $\hat{n}_w$ through θ_{mw} , carrying m onto w , and fixes every direction orthogonal to that plane. Written that way it holds on any $S^{N_{\dim}-1}$:

$$\begin{aligned} \text{transport}_{m \rightarrow w}(v) = v + (\cos \theta_{mw} - 1) [(v \cdot \hat{n}_m) \hat{n}_m + (v \cdot \hat{b}) \hat{b}] \\ + \sin \theta_{mw} [(v \cdot \hat{n}_m) \hat{b} - (v \cdot \hat{b}) \hat{n}_m], \end{aligned} \quad (35)$$

where $\hat{b} = (\hat{n}_w - (\hat{n}_m \cdot \hat{n}_w) \hat{n}_m) / \|\hat{n}_w - (\hat{n}_m \cdot \hat{n}_w) \hat{n}_m\|$ is the unit direction of w as seen from m within their plane. Unlike T , offset needs no Collapse for this: it is handed two seats, and the two seats name the turning plane themselves. At $N_{\dim} = 3$, Equation (35) is Equation (33) exactly, and the equator example below lands on the same point. In four dimensions, with $\hat{n}_m = (1, 0, 0, 0)$, $\hat{n}_k = (\frac{\sqrt{3}}{2}, 0, 0, \frac{1}{2})$, and $\hat{n}_w = (0, 1, 0, 0)$, the step $v = \log_{\hat{n}_m}(\hat{n}_k) = (0, 0, 0, \frac{\pi}{6})$ lies outside the turning plane, is carried unchanged, and lands at $\hat{n}_* = (0, \frac{\sqrt{3}}{2}, 0, \frac{1}{2})$.

4.2 A check on one equator, not a definition

Place three arrows in the xy -plane:

$$\hat{n}_m = (1, 0, 0), \quad \hat{n}_k = \left(\frac{\sqrt{3}}{2}, \frac{1}{2}, 0\right), \quad \hat{n}_w = (0, 1, 0)$$

(0° , 30° , 90°). Then $\log_{\hat{n}_m}(\hat{n}_k) = (0, \pi/6, 0)$. Here $\theta_{mw} = 90^\circ$ and $\hat{u} = (0, 0, 1)$, so (33) is a quarter-turn about $\hat{z}$ and gives $(-\pi/6, 0, 0)$. The exponential at $\hat{n}_w$ lands at

$$\hat{n}_* = \left(-\frac{1}{2}, \frac{\sqrt{3}}{2}, 0\right) \quad (120^\circ).$$

On this slice the walk *looks like* $90^\circ + 30^\circ = 120^\circ$. That sum is a protractor check. It is not the operator, and it is not a night. If a lexicon writes the names Sun, Day, Moon on the first three seats, it may write Night on the landing: whatever Sun is to Day, Moon is to that fourth place. The geometry only promised a fourth pin. The same moral as weather: Day and Night are dictionary pins, not sums of luminaries. The step that carries Sun to Day is a walk, and the walk is what is copied; $\uplus$ of those seats does not land on a time of day. Whether the two halves really are the same walk depends entirely on how a lexicon seated them, which is a dictionary claim and not one this paper makes. Offset only copies a walk; a later comparison may find the landing *near* an already seated name. It does not mint the name.

4.3 Offset is not T

$T_w(x)$ spins an occupant about an axis. Offset plants a new seat from three seats. A path does not need (31). A dictionary act does not need (20). Do not encode the walk as a ghost actor whose sole job is to trigger T .

The popular ambient sum $\hat{n}_k - \hat{n}_m + \hat{n}_w$ then normalize (Recipe 1) is the embedding fluke written as geometry. It is recorded as a cheaper cousin, not as a derivation of the landing. This paper uses Recipe 2 so the copied step stays on the skin of the ball.

5 Encodings, Sorts, and Mixed Documents

The unit sphere is a single bounded *register*. Numbers, words, and colors are not three different globes; they are three encodings of a token onto the same globe. A scientific paragraph already mixes those alphabets on one page. Pattern Arithmetic may do the same, provided each loaded state carries a *sort* and the operators respect that sort.

5.1 Two primary sorts

A state $x \in \mathcal{S}$ is the typed atom of [Axiom 1](#). What the geometric 4-tuple (r, θ, ϕ, χ) *means* depends on the sort τ and the encoding used to write it:

- **Numeric sort (num).** The state represents a scalar. One toy encoding — the only numeric encoding used in this paper — stores an integer n on a marked clock as $\chi = 2\pi n/N$ (arithmetic modulo N), at a fixed direction, with $r = 1$. Another, stereo-style, is recorded as an option: it would send $x \rightarrow \pm\infty$ toward the poles so that unbounded reals still occupy the globe. It is not used below.
- **Lexical sort (word).** The state is a labeled point in a toy dictionary living on the same globe. Similarity is the angle between points (cosine similarity). A bag of words is formed by $\uplus$. A sentence or poem is formed by T .

Schoolbook addition lives only in **num** mode. MixMerge as a *domain* reading lives in the Tier-1 sorts of [Axiom 0](#): **color** under this paper’s phasor inclusion, and **note-pitch**, whose octave equivalence makes the circle native rather than borrowed. The two sorts worked below are **num** and **word**; neither takes a domain MixMerge. Two **word** states may be mixed as geometry; that is not a χ -reading of language. The two equations may share a sphere. They may not share an operator.

5.2 Numeric mode: $7 + 5 = 12$ on the clock, and why $\uplus$ is not 12

Fix the modulus $M = 24$ and the encoding

$$E_{\text{num}}(k) = (r = 1, \hat{n} = (1, 0, 0), \chi = 2\pi k/24, \tau = \text{num}), \quad k \in \mathbb{Z}/24\mathbb{Z}. \quad (36)$$

Decode is the inverse chart: $k = M\chi/(2\pi)$, an integer hour mark. Then 7 sits at $\chi = 105^\circ$, 5 at $\chi = 75^\circ$, and 12 at $\chi = 180^\circ$. Schoolbook $+$ is decoded arithmetic:

$$7 + 5 = E_{\text{num}}^{-1}(E_{\text{num}}(7)) + E_{\text{num}}^{-1}(E_{\text{num}}(5)) = 12, \quad E_{\text{num}}(12) = (1, (1, 0, 0), 180^\circ, \text{num}). \quad (37)$$

That is $7 + 5 = 12$ as a calculation on $\mathcal{S}$: decode, add in $\mathbb{Z}$, encode. Bundle is legal on the same pair but answers a different question: $7 \uplus 5 = \{7, 5\}$ records which numerals are present and is not 12. If one instead applies MixMerge (Equation (9)) to the same two geometric marks, the phasors are

$$z = e^{i7\pi/12} + e^{i5\pi/12} = 2 \cos(\pi/12) e^{i\pi/2}, \quad (38)$$

so $r' = \min(1, 2 \cos 15^\circ) = 1$ and $\chi' = 90^\circ$. The clock would read that mark as hour 6, not hour 12. A bowl $\{E_{\text{num}}(7), E_{\text{num}}(5)\}$ is two seats, still not 12. This single contrast is why the type table exists: $+$ answers *how many*; $\uplus$ and MixMerge do not.

5.3 A document is a sequence, not a blend

Words and numbers coexist in reports because a document is an ordered tape of tokens, not one mixed arrow. Write a paragraph as a sequence of typed atoms

$$(x_1, \dots, x_m), \quad x_i = (r_i, \hat{n}_i, \chi_i, \tau_i), \quad \tau_i \in \{\text{num}, \text{word}\}, \quad (39)$$

The sphere holds one current occupant of the register; the paragraph is the history of loads and actions. A poem (Section 3.5) is one species of that history: a T -walk on one occupant. Juxtaposition is another. MixMerge is not.

They may meet in only three honest ways.

Juxtaposition. The default. Tokens sit in order and no algebra is applied between them. In “the sample reached 37.4°C at dawn,” the number 37.4 is not blended with the word dawn.

Typed action. The operator $T_w(x)$ is allowed across sorts only with an explicit output sort. Two useful conventions, which the rest of this paper adopts, are:

- a word acting on a number yields a number-in-context (still **num**);
- a number acting on a word yields a refined word (still **word**).

A number may also set a numeric dial of T — for example the rotation angle $\alpha(r)$ — because that dial is already a scalar. That is how “wind of 12 m/s” can strengthen a turn without becoming a synonym of wind.

Cast. Rare and named. The map $E_{\text{num} \rightarrow \text{word}}(2) = \text{two}$ turns a digit into a spoken numeral; the inverse is parsing. Casts are dictionary maps, not $\uplus$.

The illegal meeting is an untyped blend

$$\text{dawn} \uplus 37.4, \quad (40)$$

which has no meaning until one side is first cast into the other's sort.

5.4 Operator legality

Table 2 records which buttons accept which pairs. Here $+$ stands for schoolbook arithmetic on the numeric encoding (decode, operate in $\mathbb{Z}$, encode; $+$ is the worked case of Section 5.2), $\uplus$ is idempotent Bundle, $\boxplus$ is MixMerge, T is typed action, and $x - a + b$ is the analogical offset of Section 4.

5.5 Worked mixed sentence

Take the report line *Wind of 12 m/s darkened the sun.* The register is used as follows.

1. Load wind, sort **word**.
2. Load 12, sort **num**. Juxtapose; do not form $\text{wind} \uplus 12$.
3. Load m/s as a unit-name (lexical qualifier of the number).
4. Load sun, sort **word**.
5. Let the number cap a strength, $\sigma(12) \in [0, 1]$, and set $r_{\text{wind}} = \sigma(12)$ so that $\alpha = \pi\sigma(12)$ in Equation (20).
6. Apply $T_{\text{wind}}(\text{sun})$. Output sort: **word**.

The scalar never became a hue. It turned a dial of an action whose actors remained words. That is the same division of labor a printed paper already uses: digits set magnitudes; words name what the magnitudes belong to.

5.6 Several lexicons, one concept sphere

Surface tokens are not meanings. Two dictionaries may write the same pair of concepts with unrelated strings:

$$\begin{aligned} L_1 : \quad & \text{xyz} \mapsto \text{man}, \quad \text{qyz} \mapsto \text{woman}, \\ L_2 : \quad & \text{fgh} \mapsto \text{man}, \quad \text{hgf} \mapsto \text{woman}. \end{aligned} \quad (41)$$

In the usual pipeline each lexicon is tokenized, encoded, and *embedded separately*. The geometry that makes man near woman is then learned twice, and a further alignment step is required before L_1 and L_2 can share an analogy. The token-to-embedding map is doing two jobs at once: naming a string, and inventing a place for its meaning.

The unit sphere is offered as a *proposed common coordinate register*, not as a demonstrated universal semantic space and not as a replacement for a lexicon table. A token is a surface form. A state on $\mathcal{S}$ is a meaning. Each lexicon is only a cast

$$E_L(\text{token}) = x \in \mathcal{S}, \quad (42)$$

the same kind of map as $E_{\text{num} \rightarrow \text{word}}$ above. Under that discipline both **xyz** and **fgh** load the *same* state **man**. The offset **Sun** $\rightarrow$ **Day** of Section 4 is then computed once, on the globe, and reused by every lexicon that points at those two points.

Embeddings therefore stay. They do not disappear. What changes is that they are no longer private vector spaces, one per lexicon. They are embeddings *into the same unit* $\mathcal{S}$. Once $E_{L_1}(\text{xyz}) = E_{L_2}(\text{fgh}) = \text{man}$, the two surface forms occupy one address, and every angle measured from that address is the same angle:

$$\angle(\text{man}, \text{woman}) = \angle(E_{L_1}(\text{xyz}), E_{L_1}(\text{qyz})) = \angle(E_{L_2}(\text{fgh}), E_{L_2}(\text{hgf})). \quad (43)$$

Under that proposal the embeddings would be uniform: same globe, same points for the same concepts, same angles for the same relations. A raw string still needs the lexicon map E_L . It would no longer need a new geometry. The algebra does not mint a universal **man** point by itself. That is a dictionary programme, not an axiom.

5.7 Deferred programme: standardized concept dictionaries

The present section only *names* the shared globe. It does not build the dictionaries, and it does not treat them as part of the algebra closed in Axioms 0–7. The later refinement, if the idea is pursued, is a pair of catalogues rather than a new operator, and is filed under Future Work (Section 11):

- one *concept dictionary* — a stipulated inventory of unit states on $\mathcal{S}$ (or a higher-dimensional analogue), so that **man** and **woman** have standard addresses and therefore standard angles;
- one *language lexicon* E_L per written or spoken system, sending that language’s tokens onto those addresses;
- learnable paths — the species of T and the dials inside a given T (angle map, phase map, paint-versus-poem gate) may be fitted from examples. Learning would adjust the path, not the rule that a bowl is collected before it is read. See also Section 11.

Standardization of the first catalogue is what makes embeddings uniform across English, a cipher, or a second natural language. The second catalogue is still a tokenizer-plus-encoding: it is the phone book, not the city. Tokens and encodings have already earned their keep, so they stay. They only have to work *lighter*: look up a standard address instead of inventing one. They disappear as *inventors of geometry*. They do not disappear as *names*. A world in which every language pointed at the same concept dictionary would still print letters on a page; It would no longer need each corpus to learn a private **man**–**woman** angle. That programme is recorded here and left open. It is not an axiom, and it is not implied by the existence of $\mathcal{S}$ alone.

5.8 Scope of the register model

This section claims only that one sphere can *host* both alphabets as a workspace, and that several lexicons can point at the same workspace through different casts. It does not claim that the two-dimensional surface S^2 is a realistic embedding space for a full natural-language lexicon: a working word embedding typically needs tens or hundreds of dimensions. The construction here is a typed register and a toy dictionary on one

globe. Scaling the lexical encoding is a later choice — a higher-dimensional sphere, a spherical-harmonic field, or an external lookup table that writes states into the register. The lookup tables E_L may still be learned; the claim is only that they should write into one geometry rather than each minting a private one. The unit ball of radius 1 in $\mathbb{R}^3$ is where the idea is taught. The unit sphere in the embedding dimension is where close semantics can live. The algebra — Bundle, path, result — is meant to travel with the dimension. It is not meant to replace 768 numbers with two angles and hope the library still sorts.

6 A Sketch of Conventions

These notes are a seed for later grammars, not a code to enforce. They show one way a typed register *might* be spoken. A later grammar — learned, cultural, or formal — may replace every line. This paper only needs the ground level enough that such a grammar has somewhere to stand.

Case does not move the address. The tokens **man**, **Man**, and **MAN** load the same direction (θ, ϕ) . Orthography is a doorbell, not a new city block.

Language does not move the address. English **man** and a Polish token for the same concept load the same point. Only the lexicon map E_L changes. Language is a phone-book tag L , not a second angle.

Color is reserved for mood, not for language. Phase χ is already used as affect in the poem example (Section 3.6). If language or case also occupied χ , that example would break. Emphasis may tint χ or raise r ; it should not relocate Sun to Moon.

A bag is not a sentence. $\{\text{sun}, \text{wind}\}$ is $\uplus$. “Wind darkens the sun” is T . Ingredients and order are different verbs. How often a unit is used is not a second shell in the bag: the path may visit **twinkle** twice; **num** may count $1 + 1 + 1 + 1 = 4$. $\uplus$ will not.

A numeral is not a word until it is cast. The digit 2 stays **num**. The word **two** is **word**. They may share a bowl — Bundle is sort-blind — but $2 \uplus \text{two} = \{2, \text{two}\}$ is two seats, not a cast and not a sum. Becoming a word requires the explicit $E_{\text{num} \rightarrow \text{word}}$ of Section 5, which no Bundle performs.

A toy row makes the first three conventions visible. Place **man** at $(\theta, \phi) = (10^\circ, 40^\circ)$:

Token	L	(θ, ϕ)	χ (mood only)
man	English	$(10^\circ, 40^\circ)$	0° (neutral)
Man	English	$(10^\circ, 40^\circ)$	0°
MAN	English	$(10^\circ, 40^\circ)$	20° (emphasis, optional)
man	Polish	$(10^\circ, 40^\circ)$	0°

Same block. Different doorbells. Nothing here is binding.

7 Operational Algebra & Numerical Demonstration

7.1 MixMerge of two shells

Collecting $r = 0.6$ at 0° with $r = 0.8$ at 120° (same place) yields two shells, not an angle. Applying MixMerge ($\boxplus$) and Equation (9) yields the 4-tuple $(r', \theta', \phi', \chi') \approx (0.721, \theta, \phi, 73.9^\circ)$. The older figure 66.6° is withdrawn: Axiom 5 as a length-only formula never produced it.

7.2 Rotational transposition, and what division is not

Rotation is the rule of §2.10: $\mathbf{R} \in \text{SO}(3)$ turns $\hat{n}$. Inverse rotation is well-defined. Division as “deconvolution and signal unmixing” is not defined in this paper — shown here to be ill-posed for this algebra specifically, not merely unbuilt.

Given only a Mixed result, the bowl that produced it is not merely hard to recover; it is not determined by the result at all. The introduction’s own red-and-green example, 0° and 120° at equal strength reaching yellow at 60° , is one point in an infinite family: two components at 10° and 100° , correctly weighted, reach the identical phasor sum, to machine precision — verified directly, not merely argued, and neither pair is preferred by the result over the other. Assuming a specific composition and solving for it does not rescue this. A wrong assumed composition still reconstructs the observed result exactly: solving for 10° -and- 100° strengths against the true red-and-green result returns a residual of 10^{-17} , indistinguishable by any check internal to the algebra from the correct pair’s own residual of zero. An algebra that states only forced, verified identities has no honest place for an operation whose wrong answer looks exactly like its right one.

This is deliberately left to inference (companion paper, *Pattern Arithmetic: Inference from Bowls and Paths*), where a comparison against a witnessed reference is held to a stated confidence and may abstain — a discipline this algebra’s own MixMerge does not have available to it, and should not manufacture.

7.3 Spectral storage (optional, later)

Spherical-harmonic arrays of size $O(L^2)$ may later store a field of many atoms. They are not the working atom of Axiom 1, and $O(N_{\text{vox}}^3) \rightarrow O(L^2)$ is not a theorem of Pattern Arithmetic as it stands.

8 Higher-order systems as bowls and paths

These examples illustrate the $\boxplus$ -versus- T split with borrowed vocabulary. They are costumes (Section 1.4), not models of weather, ecology, or disease. Higher-order truths are bowls of basic units read by a path. A storm, a forest, or a cancellation is such a reading. It is an example of the spine, not a model of the sky, the woods, or the clinic. The path defines the higher-order system: the bowl only lists who is present.

A vertical pinball table makes the same split visible. Each bumper is one seated unit — a direction and a size r . The set of bumpers on the board is the bowl (wind, temperature, humidity, pressure). The pockets at the bottom are dictionary pins already named storm and calm. Dropping a ball from a slot is a path T_{weather} : the ball meets some bumpers, in some order, and each kick depends on that bumper’s angle and strength.

The landing pocket is a comparison after the walk. Adding the bumpers together does not *make* a pocket. A different slot, or a different order of kicks, is a different weather. This is one rough picture of *a* transform path, not the mechanics of every T , and not a second axiom. A MixMerge reading ($\boxplus$), and a poem path T_{poem} , need not bounce.

8.1 Atmospheric storm (costume): one pipeline with numbers

This is a costume for the spine, not a model of the sky. These walks are not weather equations, and they are not forest equations. They are still computations: ordered readings of many co-present units that can land near a higher-order name already seated in a dictionary. Four incoming states, one bowl, two paths, two endpoints. Every figure below is Rodrigues' rotation of $\hat{n}$ by $\alpha = \pi r$ about the actor's axis, with χ added. A reader can recompute.

1. Input states.

$$\begin{aligned} \text{Wind} : r = 1, \hat{n} &= (1, 0, 0), \chi = 0^\circ, \tau = \text{word}, \alpha = 180^\circ, \\ \text{Heat} : r = \frac{1}{2}, \hat{n} &= (0, 1, 0), \chi = 90^\circ, \tau = \text{word}, \alpha = 90^\circ, \\ \text{Moisture} : r = \frac{2}{3}, \hat{n} &= (0, 0, 1), \chi = 45^\circ, \tau = \text{word}, \alpha = 120^\circ, \\ \text{Eye} : r = 1, \hat{n} &= \left(\frac{\sqrt{2}}{2}, 0, \frac{\sqrt{2}}{2}\right), \chi = 0^\circ, \tau = \text{word}. \end{aligned}$$

The eye sits at $(\theta, \phi) = (45^\circ, 0^\circ)$.

2. Bundle into the bowl.

$$R = \text{Wind} \uplus \text{Heat} \uplus \text{Moisture} \uplus \text{Eye} = \{\text{Wind}, \text{Heat}, \text{Moisture}, \text{Eye}\}. \quad (44)$$

Four distinct names, four seats. A second Wind would not open a fifth. Call this bowl $s(\text{weather})$ if a name helps: it is still four pins, not a fifth vector. MixMerge is a different reading and is not the storm. In the dictionary, storm and calm are *other* seated units — separate doorbells, not sums of wind, temperature, humidity, and pressure. Mere $\uplus$ cannot reach those pins. T_{weather} is the path family that reads $s(\text{weather})$; T_{deepen} and T_{clear} below are two such walks. An endpoint may fall *near* the storm pin or the calm pin. Nearness is comparison after the walk, not addition of the parts.

3. Two trajectory paths. The register occupant is the eye. The other three shells act, in order:

$$T_{\text{deepen}} = T_{\text{Wind}} \circ T_{\text{Heat}} \circ T_{\text{Moisture}}, \quad T_{\text{clear}} = T_{\text{Moisture}} \circ T_{\text{Heat}} \circ T_{\text{Wind}}. \quad (45)$$

Deepen — moisture, then heat, then wind:

$$\begin{aligned} T_{\text{Moisture}}(\text{Eye}) &= \left(1, \left(-\frac{\sqrt{2}}{4}, \frac{\sqrt{6}}{4}, \frac{\sqrt{2}}{2}\right), 45^\circ\right), \\ T_{\text{Heat}}(\cdots) &= \left(1, \left(\frac{\sqrt{2}}{2}, \frac{\sqrt{6}}{4}, \frac{\sqrt{2}}{4}\right), 135^\circ\right), \\ T_{\text{Wind}}(\cdots) &= \left(1, \left(\frac{\sqrt{2}}{2}, -\frac{\sqrt{6}}{4}, -\frac{\sqrt{2}}{4}\right), 135^\circ\right). \end{aligned}$$

Clear — wind, then heat, then moisture:

$$\begin{aligned} T_{\text{Wind}}(\text{Eye}) &= \left(1, \left(\frac{\sqrt{2}}{2}, 0, -\frac{\sqrt{2}}{2}\right), 0^\circ\right), \\ T_{\text{Heat}}(\cdots) &= \left(1, \left(-\frac{\sqrt{2}}{2}, 0, -\frac{\sqrt{2}}{2}\right), 90^\circ\right), \\ T_{\text{Moisture}}(\cdots) &= \left(1, \left(\frac{\sqrt{2}}{4}, -\frac{\sqrt{6}}{4}, -\frac{\sqrt{2}}{2}\right), 135^\circ\right). \end{aligned}$$

4. Endpoints. T leaves the occupant’s strength at $r = 1$. Phase addition commutes, so both paths end at $\chi = 135^\circ$. The directions do not agree:

$$T_{\text{deepen}}(\text{Eye}) : \quad \hat{n} = \left(\frac{\sqrt{2}}{2}, -\frac{\sqrt{6}}{4}, -\frac{\sqrt{2}}{4}\right) \approx (0.7071, -0.6124, -0.3536), \\ (\theta, \phi) \approx (110.7^\circ, 319.1^\circ);$$

$$T_{\text{clear}}(\text{Eye}) : \quad \hat{n} = \left(\frac{\sqrt{2}}{4}, -\frac{\sqrt{6}}{4}, -\frac{\sqrt{2}}{2}\right) \approx (0.3536, -0.6124, -0.7071), \\ (\theta, \phi) = (135^\circ, 300^\circ).$$

Same bowl, two weathers. The storm is the walk. It is not a fifth pin, and it is not a claim about vorticity or shear.

8.2 Table and chair (costume)

A table is one seated unit; a chair is another. $\{\text{table}, \text{chair}\}$ is their bowl under $\uplus$ — two names present, order-blind, no claim about arrangement. A room is not read from that bowl. It is read from a path: $T_{\text{chair}}(\text{table})$, the chair acting on the table, is a placement — pulled up to the table’s edge. $T_{\text{table}}(\text{chair})$, the same two shells with the actor swapped, is a different placement — the table built around a fixed chair. $\uplus$ cannot tell these apart; it only knows both are present. T can, because T is not symmetric in its two arguments the way $\uplus$ is. As with the storm’s pockets, a named room — “dining nook,” “reading corner,” “cluttered storage” — is a separate dictionary pin, not a sum of the furniture. An endpoint may land near one pin or another; nearness is comparison after the walk, exactly as in Section 8.1. This is a costume for the same split table and chair make familiar without any physics: two objects merely being present is one reality (grocery counting: one bag of grain plus one bag of grain is two bags — schoolbook $+$, not $\uplus$); two objects placed in a specific relation is a different reality, and no amount of $\uplus$ on the parts reaches it. It is not a model of furniture, interior design, or how rooms are actually named.

The furniture case is also where the motivating hypothesis behind this whole paper is easiest to say plainly: a path T_{chair} that keeps “chair acts on table” as one typed, invertible step is a shorter thing to learn than the same fact buried inside a flat embedding, where the model must first discover, from undifferentiated vectors, that arrangement is a different kind of fact than mere co-occurrence. Keeping presence ($\uplus$) and path (T) as separate primitives is a bet that higher-order truths become easier to learn because the distinction is built in rather than re-derived from data. That bet is stated here as a hypothesis, not a result: no example in this paper demonstrates it, and one is deliberately not attempted. Section 11.1 leaves it open.

8.3 Forest (costume)

Flora, fauna, soil, moisture sit as units. $\uplus$ gathers them into one bowl — four names present, order-blind (Axiom 4). There is no $A + B = AB$. A path may read competition or succession. The forest-as-system is that reading.

8.4 Pathogen cancellation (costume)

A mark and an equal-strength anti-phase mark Bundle as two seats in one bowl (Axiom 4); $\uplus$ does not cancel them. That bowl may then be *read* by MixMerge into black (Axiom 5). That is geometry. It is not medicine.

9 The Enclosure: Seats, Rest, and Camps

Every operator before [Axiom 6](#) works entirely on the bounded stage: units go in, units or states come out, and nothing asks where anything sits. Assembly asks. This section supplies the room it asks about, and then follows what happens when units are left alone in it.

The enclosure is a costume borrowed briefly and then dropped. Many small things, left alone, falling into a few quiet piles is a familiar picture from astronomy; the units here are not stars, carry no mass, and move only by the rules below. After this paragraph the working words are *seat*, *rest*, and *camp*.

9.1 Seats

A unit, or a bowl, may carry a *seat*: a position vector $\vec{P}$ measured from one shared origin, in ordinary unbounded Euclidean space. The seat is external to the 4-tuple of [Axiom 1](#), in the same way step metadata sits on the path rather than inside the units it walks. The unit is unchanged by having one; r is still bounded, $\hat{n}$ is still a semantic address and not a physical one. For two seated units,

$$d_{AB} = |\vec{P}_B - \vec{P}_A|, \quad \hat{r}_{AB} = \frac{\vec{P}_B - \vec{P}_A}{d_{AB}}, \quad (46)$$

which completes Equation (14). Where seats exist, $\hat{r}_{AB}$ is the separation axis; no axis derived from orientation is required or used.

9.2 Rest

Seats may move; meaning does not. $\hat{n}$ and χ do not crawl, so the lexicon stays put and no unit drifts into being a different unit. What moves is $\vec{P}$, and it moves by descending the energy of Equation (14):

$$\vec{P}_i \leftarrow \vec{P}_i + \delta F_i, \quad F_i = -\nabla_{\vec{P}_i} \sum_{j \neq i} U_{ij}. \quad (47)$$

There is no mass term and no momentum carried between steps. This is deliberate rather than an omission. A unit carries no inertia, so it cannot overshoot: an isolated pair closes distance monotonically and never exceeds its own starting separation. Weighting each step by a mass-like quantity was tested and produces oscillation rather than rest, precisely where the weighting is most extreme — a ninefold difference in strength gives a 729-fold weighting under a cubed reading, and a demonstrably non-monotonic trajectory against the clean decline of Equation (47) (the printed pair of §9.4).

A direct consequence: since $F_A = -F_B$ exactly and nothing divides the resulting step, two units of any polarity move by identical magnitude toward each other. Polarity labels a bond; it does not sort the movers. A second consequence: since every force comes in an equal and opposite pair, the plain, unweighted mean of the seats is unchanged by every step.

Rest is a freeze, not a stop. Taken literally, rest arrives only once, and only at the end. The pull between two units of nonzero strength is proportional to $d(d^2 + \varepsilon^2)^{-5/2}$,

which is positive at every separation $d > 0$: no two such units are ever exactly still while apart. Given unlimited time, every group eventually joins every other and all seats of nonzero strength meet at one point. What the enclosure offers instead is a stated freeze. A configuration is *at rest at resolution η* when

$$\max_i |F_i| < \eta, \quad (48)$$

and a configuration that has reached this threshold may be frozen and taken as the considered one — the accepted solution at that resolution. Its shape, the walk that produced it, and its camps are then what is read. The threshold η is a measurement convention, named alongside the step δ and the census cut below, and not a truth of the algebra: a finer η runs longer and may see groups merge that a coarser one froze apart. It is stated on force rather than on the size of a step for a reason. A step is δF_i , so a threshold on steps would let the choice of δ decide which configurations count as rest; a threshold on force does not.

The step δ in Equation (47) is a numerical convention and not a truth of the algebra. It sets how finely the descent is sampled, not where the freeze falls. For one isolated group this is exact: the plain mean of its seats is conserved, so it comes to rest on the plain mean of its starting seats whatever δ is used, agreeing across $\delta = 0.02, 0.01, 0.005$, and 0.002 to within 10^{-15} . Where several groups pull on one another weakly, the same four steps reach the freeze at the same time and with the same camps, and their bodies agree to within 2×10^{-7} — the ordinary discretization error of sampling the slow drift between groups (§9.4). A step chosen too coarsely relative to ε_{AB} will misbehave near close approach, in the ordinary way of any fixed-step scheme; that is a limit on the arithmetic, not on the law.

Inside a group, the merge does not stop short of a point. Two members close together pull in proportion to their separation, so a frozen group’s members sit on one seat to within a distance that shrinks with η — under 2×10^{-5} at $\eta = 10^{-5}$ in the configuration of §9.4. A frozen camp has a location and a membership. It has no internal shape.

9.3 Camps and bodies

Units left alone under Equation (47) do not always fall into one pile. The energy dies as d^{-3} , so distant units are nearly irrelevant to one another and local groups settle before they feel the far ones. Several separated groups can therefore reach the freeze of Equation (48) together, each one a *camp* — a camp at resolution η , since without a freeze they would go on to merge. The number of camps is not chosen in advance; it is whatever the law, the starting seats, and η produce.

For each camp G_k that has frozen, its *body* is

$$\vec{P}_{c,k} = \frac{\sum_{i \in G_k} r_i \vec{P}_i}{\sum_{i \in G_k} r_i}, \quad (49)$$

the strength-weighted centre of its members. Nothing computes this during motion and no unit seeks it; each unit feels only the sum of pairwise pulls from its neighbours. The body is measured afterwards, by an observer, and is a description rather than a destination.

Where units carry inner structure, the weights r_i in Equation (49) may be replaced by $|V_i|$, the magnitude of a unit’s own resultant. A unit whose inner arrows cancel still sits

in the camp but barely moves its body — the body then reports surviving configuration rather than bare presence. Both readings are legitimate; they answer different questions and should not be conflated. The difference is a property of a camp read *before* its freeze. Once frozen, a camp’s members sit on one seat to the resolution of η , so every weighting returns the same body to that resolution, and the body of an isolated camp is simply the plain mean of its starting seats.

A system of several camps is described by the *set* of bodies $\{\vec{P}_{c,1}, \vec{P}_{c,2}, \dots\}$, never by one grand average over all of them. Two camps ten units apart have no member at their midpoint, and reporting one body there would assert a camp that does not exist. This is the same honesty the bowl already practises when MixMerge returns black: where the evidence is two truths, do not mint a third that nothing occupies.

Counting camps requires a stated separation cut, or a chain of short links between members. That cut is a measurement convention and not a truth of the algebra: the camps are produced by the law, the census of them is not.

A camp is a checkpoint, not a destination. Every camp is read at some freeze η , after some elapsed time; a longer wait, a finer η , or a different δ freezes a different, though usually nearby, configuration. Call each such freeze a *checkpoint*. Two checkpoints of the same starting seats — upstream or downstream in time, at a different η , or under a different δ — need not agree exactly, and §9.4 already prints the drifts: 2×10^{-7} between step sizes, 2.4×10^{-8} between weightings, 3×10^{-5} across a weak inter-group pull. None of these is a mistake; each is what the same law gives at a different checkpoint.

Treating two checkpoints as *the same* freeze therefore needs one further named convention, an *allowance* ϵ_{allow} : checkpoints whose bodies differ by less than ϵ_{allow} are read as one, and checkpoints that differ by more are read as genuinely apart. This is not a new law. It sits beside η , δ , and the census cut as a fourth measurement convention, stated once and applied consistently, not a truth of Equation (14). Every drift printed above falls well inside a stated allowance of, say, 10^{-4} ; the storm’s Deepen and Clear endpoints of §8.1, 0.500 apart in direction (28.9°), do not fall inside any allowance a reader would state honestly. An allowance too generous erases real structure; one too strict reports a new camp every time the integrator takes a different step. Naming it is the price of being able to say two checkpoints are the same thing at all.

9.4 A replayable configuration

The claims of this section are printed against one configuration, so that any reader can replay them. How a configuration is generated — by hand, or by any library under any seed — is not part of the claim; the printed seats are. Six units in two groups forty apart:

Unit	r	$\hat{n}$	starting $\vec{P}$
1	0.5	(1, 0, 0)	(0, 0, 0)
2	0.6	(0, 1, 0)	(1, 0, 0)
3	0.7	(0, 0, 1)	(0, 1, 0)
4	0.8	(1, 0, 0)	(40, 0, 0)
5	0.4	$(\frac{1}{\sqrt{2}}, \frac{1}{\sqrt{2}}, 0)$	(41, 0, 0)
6	0.9	(0, 0, 1)	(40, 1, 0)

Under Equation (47) with the freeze $\eta = 10^{-5}$ of Equation (48), the runs at $\delta = 0.02$, 0.01, 0.005, and 0.002 take 628, 1261, 2526, and 6320 steps — the same elapsed time, 12.6,

in each — and every one freezes into the same two camps, units $\{1, 2, 3\}$ and $\{4, 5, 6\}$. The bodies land at $(0.33336, 0.33333, 0)$ and $(40.33330, 0.33333, 0)$: the plain means of each group’s starting seats, $(\frac{1}{3}, \frac{1}{3}, 0)$ and $(\frac{121}{3}, \frac{1}{3}, 0)$, each drawn about 3×10^{-5} toward the other by the weak pull across the gap. The four runs agree on both bodies to within 2×10^{-7} ; the plain mean of all six seats is unchanged throughout; and the r -weighted and plain bodies of the first camp differ by 2.4×10^{-8} . Run alone, the first group comes to rest on $(\frac{1}{3}, \frac{1}{3}, 0)$ at all four steps, agreeing to within 10^{-15} . Neither camp is still in the literal sense: the two remain forty apart only because their mutual pull fell below η before they could close the distance.

The mass-weighting comparison of §9.2 is replayed on a printed pair: $r = 0.9$ at $(0, 0, 0)$ and $r = 0.1$ at $(3, 0, 0)$, both with $\hat{n} = (1, 0, 0)$, $\delta = 0.02$, 3000 steps. Under Equation (47) the separation never increases, falling from 3 to 1.80. Dividing each step by r^3 — a 729-fold weighting between the two — makes the separation increase on 804 of the 3000 steps.

A printed configuration makes a result replayable. It does not make the result right: a replay reproduces what the law did from those seats, and whether that answer is the one a reader wanted is a separate question.

9.5 What the enclosure does not supply

Assembly never repels, so nothing in the enclosure pushes members apart; separated camps persist because pull dies with distance, not because anything resists, and only until a freeze is declared. There is no final plurality: left without a freeze, every camp eventually joins every other. There is correspondingly no core, no halo, and no sorting of members by strength — structures familiar from the borrowed picture that this law does not reproduce, and should not be graded against. A camp does not acquire a name by coming to rest. Naming is comparison against a seated pin, which is a different question asked afterwards and elsewhere.

10 Discussion & Related Work

Each paragraph below states what is reused and what is added. None of these citations is a claim of identity. The list is a boundary-drawing exercise against formalisms that share a sphere, an angle, or a rotation (Section 1.4). It is not a unification of those fields, and it is not a proposal to replace quantum mechanics, embeddings, or transformers. Empty headings are omitted.

Bloch sphere. A pure qubit is a point on S^2 ; a mixed qubit lives in the ball. Pattern Arithmetic reuses that unit globe. It differs in three respects. Strength $r \leq 1$ is an activation of one arrow, not a purity parameter of a density matrix. Phase χ is an internal colour of that arrow, not a global $U(1)$ that can be quotiented away. There is no measurement postulate: MixMerge may mint a third unit or black; $\uplus$ only collects. They do not Bundle a superposition onto an eigenstate. Related work may say: we reuse a unit globe; we do not borrow quantum authority for $\uplus$.

Word2vec and cosine embeddings. Lexical similarity as an angle is reused. A bag of words as a commutative Bundle is reused in spirit. What is added is a bounded unit

($r \leq 1$), an explicit collection operator that is idempotent ($A \uplus A = A$), and an order-dependent path T that cosine addition does not supply. Offset is a manifold cousin of the familiar king – man + woman arithmetic, not a proof that the toy globe is a full lexicon.

Transformers. A transformer occupies the same *slot* as T : it is a path through seated tokens. It is not the same device. The usual stack is

$$\text{sentence} \rightarrow \text{token} \rightarrow \text{encode} \rightarrow \text{embed} \rightarrow \text{transformer}.$$

The stack of this paper is

$$\text{sentence} \rightarrow \text{token} \rightarrow E_L \rightarrow \text{sphere} \rightarrow \uplus (\text{bowl}) + T (\text{walk}).$$

Tokens and encodings stay. The embedding is a cast into one sphere (of whatever dimension the lexicon needs). A higher-order unit (a line, a report, a lexicon) may treat each information atom as the already-furnished stand-in for a token: that is a closure over many units, not a change to the atom. T_{poem} , T_{prose} , and T_{equation} are species of walk, not extra nets. A transformer may later fill those walks as a learned path. This paper does not claim to replace attention; it splits presence from order so that a learned path need not also invent the city.

RotatE. Knowledge-graph embeddings that treat a relation as a rotation in a complex plane are the nearest neighbour of T . Pattern Arithmetic extends that idea from S^1 (or a product of circles) to $\text{SO}(3)$ acting on S^2 , and pairs the rotation with an additive phase χ . It also keeps a separate Bundle operator, which RotatE does not.

von Mises–Fisher. Directional statistics on the sphere reuse the same stage. Those models estimate a density. Pattern Arithmetic does not: the atom is a point, r is not a concentration κ , and $\uplus$, $\boxplus$, T , and offset are algebraic operations on units, not maximum-likelihood fits.

Riemann sphere. Stereographic compactification of $\mathbb{C} \cup \{\infty\}$ onto S^2 is the natural home of one numeric encoding (unbounded reals toward the poles). The algebra of this paper is not the Möbius group. A clock, or a stereo map, is a chart; $+$ in **num** mode is decoded arithmetic, not a fractional linear map.

SO(3). In three dimensions, the transformation T is a named action of the rotation group, plus a phase shift; beyond three, each T is a simple rotation — one plane turned, the rest held — and not a general element of $\text{SO}(N_{\text{dim}})$ (§3.4). That is the whole geometric engine of a path. What $\text{SO}(3)$ does not supply is the bowl, the cap $r \leq 1$, the sort τ , or MixMerge. Non-commutativity of poems is inherited; idempotence of $\uplus$ is not.

Spherical harmonics. A later field representation of many arrows could be stored as Y_ℓ^m coefficients, replacing an $O(N_{\text{vox}}^3)$ voxel cube by an $O(L^2)$ spectral array. That is a compression option for a field, not a definition of the atom, and it is not claimed as a result here.

11 Conclusion & Future Directions

The claim of this paper is narrow (Section 1.4). Counting is not the only arithmetic for combining units of information. A combination rule should match the unit. The sphere and the operators $\oplus$, $\boxplus$, T , offset, and $\boxtimes$ are one proposed encoding of that claim. They are not a general theory of language, physics, biology, or computation. The costumes, the enclosure, and related-work paragraphs are illustrations and boundaries, not domain models.

What was written is two pipelines on the same unit. Bundle many sorted units into a bowl of typed atoms $x = (r, \hat{n}, \chi, \tau)$ on a unit sphere, then derive results by ordered transformation of a register occupant. Separately, assemble units without fusing them and let external seats descend until a stated threshold freezes them; the frozen camps have bodies, not names. Naming a camp is left open here, on purpose: a second companion paper, *Pattern Arithmetic: Higher Truths from Assembly*, now states it as the question still to be answered, not a comparison already made. Embeddings, on this view, are not abolished. They are embeddings into one proposed common coordinate register, so that a concept such as man *may* keep the same angles under every lexicon that points at it. That is a design, not a demonstrated universal semantic space.

A few spoken conventions are sketched in Section 6 as seed, not statute.

11.1 Future work

What remains to be built, and is only filed here, is a standardized concept dictionary together with per-language lexica E_L (Section 5.7). Until those catalogues exist, tokenization and encoding remain the handles by which a paragraph finds its points. They may continue; they will only have to work lighter. The open bet is that they need not also remain the architects of a new space of meanings. The algebra closed in this paper does not produce a universal man point by itself. Standardization is sociological and empirical; it is not an axiom.

A further deferred question is whether a path T should stay hand-specified or carry learnable weights: the angle law $\alpha(r)$, a phase map, a gate between MixMerge ($\boxplus$) and T_{poem} , and the lexicon casts E_L . The honest constraint is structural. If undo-the-poem matters, learning should stay inside $\text{SO}(3)$, not Bundle T into an arbitrary matrix. Weights may change how the shells are thrown. They should not let the shells throw themselves.

Left fully open is the hypothesis raised in Section 8.2: that separating presence ($\oplus$) from path (T) as distinct primitives should make higher-order truths cheaper to learn than recovering the same distinction from undifferentiated vectors. Testing this would need a controlled comparison — a compositional task, a flat-embedding baseline, and a T -path model, measured on data efficiency rather than final accuracy — and no such comparison is attempted here. The table-and-chair and storm costumes motivate the hypothesis; they do not test it.

Ambient $\mathbb{R}^3$ makes the expressions $\mathcal{S} + (a, b, c)$ and $\mathcal{S} \cdot (a, b, c)$ look legal. They are legal as *vector* arithmetic. They are not automatically legal as *unit* arithmetic. A state with $r = 0.98$ added to another with $r = 0.98$ along the same direction is $r = 1.96$, which has left the ball. A scalar $2 \cdot (0.98)$ does the same. The unit cap $r \in [0, 1]$ then forces a choice, to be made later and not here: clip (overflow saturates at 1), project (renormalize direction and forget extra strength), or refuse (the operator is undefined).

off the ball). Schoolbook + on the `num` encoding is a different door: decode, add in $\mathbb{R}$, encode back if the chart allows. Do not let $\mathbb{R}^3$ addition quietly replace $\oplus$, $\boxplus$, or T . The embedding lives in $\mathbb{R}^3$. The unit lives in the ball. As already noted in Section 5.8, the three-dimensional ball teaches the algebra; close document semantics, if pursued, live on a higher-dimensional sphere that carries the same verbs.

Dimensionality. Nothing stated above is intrinsically three-dimensional; the stage may be $S^{N_{\text{dim}}-1}$ as well as S^2 . The magnitude cap $r \in [0, 1]$ is a property of a vector's length, not of the ambient dimension. The log and exponential maps used for the offset operator (Equations around (30)–(31)) are already Riemannian statements that hold on any $S^{N_{\text{dim}}-1}$ in $\mathbb{R}^{N_{\text{dim}}}$, not only S^2 . $\text{SO}(3)$ generalizes to $\text{SO}(N_{\text{dim}})$ as an abstract group — associativity and invertibility are group-theoretic facts about $\text{SO}(N_{\text{dim}})$ for any N_{dim} , not special to $N_{\text{dim}} = 3$ — but that is a statement about the group, not about T_w 's specific construction from one vector: §3.4 found that a single $\hat{n}_w$ does not carry enough data to fix a rotation outside three dimensions, however true the group's own properties remain. Three dimensions were chosen for legibility and hand-checkable worked examples, not necessity. Raising N_{dim} is expected to relieve one specific problem: too few directions forcing distinct concepts to crowd onto nearby points, a kind of clipping by dimensional scarcity rather than by magnitude. Because crowding is a many-atom problem, the benefit is conditional on the bowl actually being large enough to crowd; raising N_{dim} for a bowl of two or three atoms relieves nothing, since ordinary distinct points already sit far apart on S^2 . It does *not* relieve the separate overflow problem discussed just above (clip, project, or refuse when r exceeds 1) — that is a magnitude question and stays open regardless of N_{dim} .

T and offset are no longer part of this TODO. §3.4 restates T for general N_{dim} , invertible and non-commutative up to $N_{\text{dim}} = 50$ and exact at $N_{\text{dim}} = 3$; §4.1 restates offset's transport as a rotation within the plane of its two seats, exact at $N_{\text{dim}} = 3$. The remaining operators contain no step that depends on the dimension. Bundle is set union. MixMerge reads phase among units that already share a direction. Assembly's bracket $3 + \hat{n}_A \cdot \hat{n}_B$ stays within $[2, 4]$ in any dimension, so it never repels by the same argument as in three. The identities of Axiom 7 are dot-product identities. Each follows from the argument just given, in any dimension; no run is needed to establish it. What stays three-dimensional is presentation: the polar chart (θ, ϕ) and every worked example, by the convention stated at the start of this paper. MixMerge across differing directions remains open (§11.1), but that is a question about direction, not dimension.

The algebra is dimension-blind and count-honest: every informational arrow is a unit; the stage may be S^2 or $S^{N_{\text{dim}}-1}$; packing more arrows into one 3-D atom is a density choice, not a change of claim. 3-D with a high N is the right teaching object. Higher N_{dim} is the same object with more sky — use it when crowding, not when the idea feels small.

Open: MixMerge across differing directions. Axiom 5 defines MixMerge only among states that share a direction $\hat{n}$, where it is phasor combination on χ . What MixMerge should mean for states at different directions is not settled here. An earlier draft of this paper carried a provisional recipe — normalize the vector sum $r_A \hat{n}_A + r_B \hat{n}_B$ — and it is withdrawn, because it is the only place in the construction where $\hat{n}$ would move by something other than a rotation. Everywhere else, direction is the exclusive property of $\text{SO}(3)$ (§2.10) and of T , and a blend that slides an arrow between axes is

not a rotation. Three routes remain open, and this paper picks none of them: refuse MixMerge across directions outright, making the operator partial in place just as it is already partial in sort; rotate one state onto the other's axis by an explicit, stated T first, so that all blending stays phasor blending; or accept a genuine two-stage operator and state separately what it preserves. Whichever is chosen, it should be shown that the resulting map still respects the cap and does not smuggle $\text{SO}(3)$'s job into $\boxplus$.

Open: reducing a path by merging close actors. Two *sequential* actors already compose into one equivalent rotation for free — that is just $\text{SO}(3)$'s (or $\text{SO}(N_{\text{dim}})$'s) closure, already latent in T 's associativity. That is not the open question. The open question is whether several actors that are merely *close* — by some stated distance on $\mathcal{S}$, not identical — may be replaced by one representative actor before a path is walked, trading path length for approximation. This is a coarse-graining or clustering step, not composition, and it is not defined anywhere in this paper. Two things would need stating before it could be trusted: a distance and a merge rule (is the representative a $\boxplus$ -blend of the close actors, an average rotation, or something else), and a bound relating that distance to how far the reduced path's endpoint may drift from the unreduced path's endpoint. Composition is associative and exact; this proposed reduction is neither, and its error, if any, has not been characterized. A merge of this kind, whatever form it eventually takes, would produce an additional, derived actor for a shorter path to use — it would not remove the original close actors from the bowl. The bowl under $\boxplus$ (Section 2.2) already keeps every distinct seated unit; a reduced actor sits alongside its sources, not in place of them, and the unreduced path through the original actors remains walkable and exact. Reduction would be a second, optional reading of the same bowl, the way MixMerge and T already are, not a rewrite of what the bowl holds. Left open.

Open: location as a sort. A position (latitude and longitude, say, or any other two-angle coordinate) is not a fifth coordinate on every atom; the atom stays the 4-tuple of [Axiom 1](#), and $\hat{n}$ remains a pattern's semantic address, not its physical one. A location instead sits for admission the way `chladni` did (Section 2.9): its own τ , its own chart, and its own φ_τ stating what native law on positions $\boxplus$ or T would have to honor. Neither has been attempted here. Two candidates worth weighing when it is: whether combining two locations behaves enough like $\boxplus$'s union of presence to admit at Tier 0 only — plausible if the native operation absorbs (a smaller region inside a larger one yielding the larger), since $\boxplus$ has no absorption law and cannot honor one; or whether some native law on positions maps onto T or $\boxplus$ at Tier 1, the way `chladni`'s modal superposition did. Left open.

This is a separate question from the seat $\vec{P}$ of Section 9.1, and admitting one would not settle the other. A seat is external bookkeeping — a place a unit occupies, carrying no meaning and entering no chart. A location *sort* would be a unit in its own right, seated in bowls alongside temperature and wind, and would have to pass Axiom 0 on its own terms before $\boxplus$ or T could touch it.

Pair	+	$\uplus$	$\boxplus$	$T_w(x)$	$x - a + b$
(word, word)	no	yes	yes [†]	yes	yes
(num, num)	yes	yes [‡]	no	no*	no*
(word, num)	no	yes [‡]	no	typed	no
(num, word)	no	yes [‡]	no	typed	no
(color, color)	no	yes	yes	yes	open
(note-pitch, note-pitch)	no	yes	yes	yes	open
(chladni, chladni)	no	yes	yes [§]	open	open
(τ_1, τ_2) any other	no	yes [‡]	open	open	open

Table 2: Sort discipline on a single sphere. Every $\boxplus$ entry is additionally subject to the place restriction of [Axiom 5](#): MixMerge combines χ at a common $\hat{n}$, and MixMerge across differing directions is open (Section 11.1). **color** and **note-pitch** are the Tier-1 sorts of Table 1; they Bundle like any other sort, and their MixMerge is a domain reading rather than inert geometry. MixMerge on **(num, num)** is the clock trap of Section 5.2: it is geometrically defined and type-refused. [†]MixMerge on **(word, word)** is allowed as geometry and carries no linguistic χ -claim ([Axiom 0](#), Tier 0). *Numeric T and numeric offset may be defined later as maps of the chosen numeric encoding; they are not the lexical operators of Sections 3 and 4. Offset on the Tier-1 sorts is untested here and marked open rather than claimed. [‡]Bundle is sort-blind: the $\uplus$ column is *yes* for every pair, including pairs of different sorts and sorts not listed here. A bowl records which units are present, and presence carries no domain claim, so there is nothing for a sort mismatch to corrupt. What Bundle does *not* do is combine: $2 \uplus \text{two} = \{2, \text{two}\}$ is two seats and no cast, and $1 \uplus 1 = \{1\}$ is one seat and not the sum 2. Counting stays with + (Section 5.2). The final row is the general rule; a sort seated on the sphere may always be Bundled, while its other operators wait on [Axiom 0](#). A sort inadmissible for Tier 1, such as **snowflake** in Table 1, still falls under this general row at Tier 0: it may be Bundled, and its other operators remain open until a stated map is given — the same status Table 1 already grants it, not a stricter one. What has genuinely not been done for **snowflake**, unlike **num**, **word**, **color**, or **chladni**, is the construction of an explicit chart; Tier-0 admission is asserted here, not built. A sort with *no* row in either table — one this paper has not considered at all — is undecided in the stronger sense of Table 1’s closing line, and is not seated until one is given. [§]MixMerge on **(chladni, chladni)** is a domain reading only within one mode, where the place restriction above is satisfied automatically; across modes it is undefined rather than merely untested (Section 2.9). T is marked open because direction carries a mode label for this sort and no domain meaning is claimed for rotating it.